SIOB 296 Introduction to Programming with R

Eric Archer (eric.archer@noaa.gov)

Week 1: April 4, 2017

Reading

The Book of R:
Chapter 1, pages 3-5 (installing R)
Chapter 2 (Numeric, Arithmetic, Assignment and Vectors)
Chapter 4 (Non-Numeric Values)
Chapter 6, pages 103-114 (Some Special Values)

R Console

- commands and assignments executed or evaluated immediately
- separated by new line (Enter/Return) or semicolon
- recall commands with ↑ or ↓
- case sensitive

NB: EVERY command is executing some function and returns something

Help

There are several ways of getting help. The most common is just the `help` command:

```
help(mean)
```

This can be shortened to just `?` in most cases:

```
?median
```

For some special functions, topics, or operators, you should use quotes:

```
help("[")
```

The examples in help pages can be run using the `example` function:

```
example(mean)
```

```
mean> x <- c(0:10, 50)

mean> xm <- mean(x)

mean> c(xm, mean(x, trim = 0.10))
[1] 8.75 5.50
```

Finally, if you don't know the name of the function, but you know a keyword, you can use `help.search`:

```
help.search("regression")
```

Workspace

The contents of the workspace can be viewed with `ls`:

```
ls()
```

```
[1] "x"  "xm"
```

Useful workspace functions
`rm()`: remove an object
`rm(list = ls())`: remove all objects in the workspace
`save.image()`: save all objects in the workspace
`load(".rdata")`: load saved workspace
`history()`: view saved history
`#`: comment

Math

The R console can be used as a powerful calculator where both complex and simple calculations can be made on the fly:

```
4 + 5
```

```
[1] 9
```

```
5 / 23
```

```
[1] 0.2173913
```

```
1 / 1.6 + 1
```

```
[1] 1.625
```

```
(-5 + sqrt(5 ^ 2 - (4 * 3 * 2))) / (2 * 3)
```

```
[1] -0.6666667
```

Other common mathematical operators can be found with `?Arithmetic`.

Data structures

There are six basic storage **modes** that you will encounter in most of your R work:
`logical`: TRUE, FALSE, T, F
`integer`: whole numbers (e.g., 1, -1, 15, 0)
`double`: double precision decimals (e.g., 3.14, 1e-5, 2.0)
`character`: character strings (e.g., "Hello World", "I love R", "22.3")
`list`: A collection of objects that can be of different modes
`function`: A set of commands initiated by a call that takes arguments and returns a value

There are six basic object **classes** that you should become familiar with:
`vector`: One dimensional, all elements are of same mode
`factor`: One dimensional, categorical data represented by integers mapped to levels
`matrix`: Two dimensional, all elements are of same mode
`array`: Multi-dimensional, all elements are of same mode
`list`: One dimensional, elements can be of different modes
`data.frame`: Two dimensional, each column is an element of same length (rows)

Special Values
NULL: Empty object or object does not exist NA: Missing data
NaN: Not a Number (0/0)
Inf / -Inf: Infinity (1/0)

Object Information
`str`: Display the structure of an object
`mode`: The storage mode of an object
`class`: The class of an object
`is.<class>`: Test if an object is of a given class

Vectors

Objects are assigned values using the "left arrow" (`<-`) operator, like this:

```r
x <- 1
x
```

```
[1] 1
```

You can also use `=` for assignment, but I seriously recommend not getting into the habit of doing that. It can actually make code harder to read because `=` is used in a slightly different context. I have found it better to be consistent and stick with `<-`.

```r
# The ':' operator creates a numeric vector incrementing by 1
x <- 1:10
x
```

```
 [1]  1  2  3  4  5  6  7  8  9 10
```

```r
# The `c` function creates a vector containing the arguments inside
y <- c("a", "b", "d")
y
```

```
[1] "a" "b" "d"
```

```r
str(x)
```

```
 int [1:10] 1 2 3 4 5 6 7 8 9 10
```

```r
is.numeric(x)
```

```
[1] TRUE
```

```r
class(y)
```

```
[1] "character"
```

```r
mode(x)
```

```
[1] "numeric"
```

Indexing

There are three ways to index any object in R:
Numeric: Using integers to reference the element number
Character: If the object has "names", using characters to specify those names
Logical: Return only the elements that match `TRUE` values

Numeric Indexing

```r
x <- 21:30
x
```

```
 [1] 21 22 23 24 25 26 27 28 29 30
```

```r
# The fifth element
x[5]
```

```
[1] 25
```

```r
# The first three elements
x[1:3]
```

```
[1] 21 22 23
```

```r
# The first, fifth, and sixth elements
x[c(1, 5, 6)]
```

```
[1] 21 25 26
```

```r
# Numerical indexing returns elements in the order they were requested
x[c(8, 9, 3)]
```

```
[1] 28 29 23
```

```r
# Replication of elements is allowed and will be acommodated
x[c(4, 6, 5, 6, 4)]
```

```
[1] 24 26 25 26 24
```

```r
# Any numeric vector is allowed
x[c(1:4, 5, 10:8)]
```

```
[1] 21 22 23 24 25 30 29 28
```

```r
# Negative numbers return all elements except the negative value
x[-3]
```

```
[1] 21 22 24 25 26 27 28 29 30
```

```r
# Don't fall into this trap
x[-1:5]
```

```
Error in x[-1:5]: only 0's may be mixed with negative subscripts
```

```r
# What you probably mean is this
x[-(1:5)]
```

```
[1] 26 27 28 29 30
```

Character Indexing

To use character indexing, you have to provide **names** to the vector

```r
names(x) <- letters[1:10]
x
```

```
 a  b  c  d  e  f  g  h  i  j
21 22 23 24 25 26 27 28 29 30
```

```r
str(x)
```

```
 Named int [1:10] 21 22 23 24 25 26 27 28 29 30
 - attr(*, "names")= chr [1:10] "a" "b" "c" "d" ...
```

Then, elements can be specified by name

```r
x["d"]
```

```
 d
24
```

```r
x[c("f", "a")]
```

```
 f  a
26 21
```

Specific names can be changed by referencing the `names(x)` vector

```r
names(x)[4] <- "fourth"
x["fourth"]
```

```
fourth
    24
```

Logical indexing

The third way to index is using logical vectors. Only elements matching `TRUE` values are returned

```r
y <- 1:4
y[c(T, T, F, T)]
```

```
[1] 1 2 4
```

Here are the primary logical operators:
`!` : Not - negates the value (!T = F, !F = T)
`&` : And - Result is T if both values are T (T & T = T, T & F = F, F & F = F)
`|` : Or - Result is T if one value is T (T | T = T, T | F = T, F | F = F)
`<, >` : Less, greater than
`<=, >=` : Less than or equal to, greater than or equal to
`==` : Equal to
`!=` : Not equal to
`any()` : Returns T if any value is T
`all()` : Returns T if all values are T

```r
x <- 50:20
x
```

```
 [1] 50 49 48 47 46 45 44 43 42 41 40 39 38 37 36 35 34 33 32 31 30 29 28
[24] 27 26 25 24 23 22 21 20
```

```r
x[x < 30]
```

```
 [1] 29 28 27 26 25 24 23 22 21 20
```

```r
x[x < 40 & x > 25]
```

```
 [1] 39 38 37 36 35 34 33 32 31 30 29 28 27 26
```

```r
x[x < 25 | x > 43]
```

```
[1] 50 49 48 47 46 45 44 24 23 22 21 20
```

Vectorization

A key component of R operations on vectors is the idea of "vectorization". In essence, this means that operations between multiple R vectors will tend to recycle elements in the smaller object to the size of the larger object. This is most easily seen in vector algebra:

```
# Add two vectors of equal length
1:5 + 21:25
```

```
[1] 22 24 26 28 30
```

```
# Add two vectors where one is a multiple of the other
1:10 + 1:2
```

```
 [1]  2  4  4  6  6  8  8 10 10 12
```

```
# Add two vectors where one is not the multiple of the other
1:10 + 1:3
```

```
Warning in 1:10 + 1:3: longer object length is not a multiple of shorter
object length
```

```
 [1]  2  4  6  5  7  9  8 10 12 11
```

Vectorization can be used in logical indexing too

```
# Select every other element
x <- 1:10
x[c(T, F)]
```

```
[1] 1 3 5 7 9
```

```
# Select every third element
x[c(T, F, F, F)]
```

```
[1] 1 5 9
```

Character vectors

```
x <- c("A", "b", "C")
x[2]
```

```
[1] "b"
```

```
# Add names with vector
y <- 1:3
names(y) <- x

# Select using logical
x[x == "C"]
```

```
[1] "C"
```

```
x[x != "b"]
```

```
[1] "A" "C"
```

```r
# Index one vector with another
x[y == 2]
```

```
[1] "b"
```

```r
# Two special values that provide a vector of lower and upper case letters:
letters
```

```
 [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q"
[18] "r" "s" "t" "u" "v" "w" "x" "y" "z"
```

```r
LETTERS
```

```
 [1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q"
[18] "R" "S" "T" "U" "V" "W" "X" "Y" "Z"
```

Logical vectors

```r
x <- c(T, F, T, F, F, T)
any(x)
```

```
[1] TRUE
```

```r
all(x)
```

```
[1] FALSE
```

```r
# Negate the vector
!x
```

```
[1] FALSE  TRUE FALSE  TRUE  TRUE FALSE
```

```r
# Every other value
x[c(F, T)]
```

```
[1] FALSE FALSE  TRUE
```

```r
# Just the TRUE values
x[x]
```

```
[1] TRUE TRUE TRUE
```

```r
# Logical vectorization
x & T
```

```
[1]  TRUE FALSE  TRUE FALSE FALSE  TRUE
```

```r
x | c(F, T)
```

```
[1]  TRUE  TRUE  TRUE  TRUE FALSE  TRUE
```

Factors

Factors are special vectors where the unique values are stored as numbers and mapped to character levels

```r
x <- factor(c("yellow", "blue", "green", "blue", "Blue", "yellow"))
x
```

```
[1] yellow blue   green  blue   Blue   yellow
Levels: blue Blue green yellow
```

```r
# Notice that the values are numerics
str(x)
```

```
 Factor w/ 4 levels "blue","Blue",..: 4 1 3 1 2 4
```

```r
# ... but the class isn't
is.numeric(x)
```

```
[1] FALSE
```

```r
# ... nor is it character
is.character(x)
```

```
[1] FALSE
```

```r
# Here's the class
class(x)
```

```
[1] "factor"
```

```r
# and the storage mode
mode(x)
```

```
[1] "numeric"
```

The numeric and original character vectors can be obtained by **coercion** using the `as.<class>` set of functions:

```r
as.numeric(x)
```

```
[1] 4 1 3 1 2 4
```

```r
as.character(x)
```

```
[1] "yellow" "blue"    "green"  "blue"    "Blue"    "yellow"
```

A factor has both `levels` and `labels`. The `levels` are the set of values that might have existed in the original vector and the `labels` are the representations of the `levels`.

```r
# The sample function takes a random sample from a vector with or without replacement
x <- sample(x = letters[1:4], size = 10, replace = TRUE)
xf <- factor(x)
xf
```

```
 [1] b b a d b c b c c b
Levels: a b c d
```

```r
# Here are the levels
levels(xf)
```

```
[1] "a" "b" "c" "d"
```

```r
# We can change the order of the levels (note doesn't change order of values in vector)
xf.lvl <- factor(x, levels = c("c", "b", "d", "a"))
xf.lvl
```

```
 [1] b b a d b c b c c b
Levels: c b d a
```

```r
# Adding a level that doesn't exist has no effect on data, but includes level in list of levels
xf.lvl <- factor(x, levels = c("c", "e", "b", "d", "a"))
xf.lvl
```

```
 [1] b b a d b c b c c b
```

```
Levels: c e b d a
```

```r
# Omitting a level causes all values with that level to be NA
xf.lvl <- factor(x, levels = c("b", "d", "a"))
xf.lvl
```

```
 [1] b    b    a    d    b    <NA> b    <NA> <NA> b
Levels: b d a
```

```r
# Labels will match order of levels
xf.lbl <- factor(x, labels = c("Z", "Y", "X", "W"))
xf.lbl
```

```
 [1] Y Y Z W Y X Y X X Y
Levels: Z Y X W
```

```r
# But you must have as many labels as levels
xf.lbl <- factor(x, labels = c("Z", "Y", "X"))
```

```
Error in factor(x, labels = c("Z", "Y", "X")): invalid 'labels'; length 3 should be 1 or 4
```