

SIOB 296 Introduction to Programming with R

Eric Archer (eric.archer@noaa.gov)

Week 2: April 11, 2017

Reading

The Book of R:

Chapter 3 Matrices and Arrays

Chapter 5 Lists and Data Frames

Chapter 6.2.4 As-Dot Conversion Functions

The Art of R:

Chapter 3 Matrices and Arrays Chapter 4 Lists Chapter 5 Data Frames

sample

The `sample` function is a very useful function for drawing random samples from vectors either with or without replacement.

If it is called with just a vector, it returns a permuted form of that vector:

```
sample(1:10)
```

```
[1] 4 7 9 1 6 2 8 10 5 3
```

A smaller vector can be created by specifying the `size` argument:

```
sample(1:10, size = 5)
```

```
[1] 8 10 5 7 6
```

In this case, 5 unique values are returned. If you want to sample with replacement, specify `replace = TRUE`:

```
sample(1:10, size = 5, replace = TRUE)
```

```
[1] 7 9 5 9 9
```

If a larger vector is to be sampled, then `replace` has to be `TRUE`:

```
sample(1:10, size = 100, replace = TRUE)
```

```
[1] 3 1 5 7 8 3 6 7 9 7 8 3 2 7 7 4 3 6 8 9 9 7 10
[24] 4 1 7 7 8 9 4 5 2 2 3 4 10 4 9 4 3 10 4 10 3 8 7
[47] 7 3 7 9 4 5 9 2 7 1 5 4 6 8 2 9 6 10 6 3 10 10 10
[70] 4 3 1 9 5 2 3 8 6 9 10 1 1 10 9 1 1 3 8 3 9 3 1
[93] 10 3 8 8 9 5 4 7
```

By default, all elements in the vector have the same likelihood of being sampled. Weights can be applied by specifying them in the `prob` argument:

```
sample(letters[1:5], size = 20, replace = TRUE, prob = c(10, 10, 1, 1, 0))
```

```
[1] "a" "a" "b" "a" "b" "a" "a" "a" "b" "c" "b" "a" "b" "c" "a" "b" "b"
[18] "b" "d" "a"
```

Matrices

Matrices are always two-dimensional objects having a certain number of rows and columns. They contain only one kind (atomic mode) of data (e.g., numeric, character, logical). They are created by supplying a vector of values to the `matrix()` function and specifying how many rows and/or how many columns to dimension it by

```
# Create a matrix
```

```
x <- 1:24  
mat <- matrix(x, nrow = 4)  
mat
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]  
[1,]    1    5    9   13   17   21  
[2,]    2    6   10   14   18   22  
[3,]    3    7   11   15   19   23  
[4,]    4    8   12   16   20   24
```

```
# How many elements are in the matrix?
```

```
length(mat)
```

```
[1] 24
```

```
#How many rows and columns?
```

```
nrow(mat)
```

```
[1] 4
```

```
ncol(mat)
```

```
[1] 6
```

Cells are selected by [row, column]

```
mat[2, 3]
```

```
[1] 10
```

Selecting a single row or single column returns a vector

```
mat[3, ]
```

```
[1]  3  7 11 15 19 23
```

```
mat[, 4]
```

```
[1] 13 14 15 16
```

Use `drop = F` to select a single row or column and return a matrix

```
mat[4, , drop = F]
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]  
[1,]    4    8   12   16   20   24
```

```
mat[, 2, drop = F]
```

```
      [,1]  
[1,]    5  
[2,]    6  
[3,]    7  
[4,]    8
```

Select several rows or columns

```
mat[c(1, 3, 4), ]
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]     1     5     9    13    17    21
[2,]     3     7    11    15    19    23
[3,]     4     8    12    16    20    24
```

```
mat[, 2:5]
```

```
      [,1] [,2] [,3] [,4]
[1,]     5     9    13    17
[2,]     6    10    14    18
[3,]     7    11    15    19
[4,]     8    12    16    20
```

Select rows, exclude columns

```
mat[1:3, -(2:4)]
```

```
      [,1] [,2] [,3]
[1,]     1    17    21
[2,]     2    18    22
[3,]     3    19    23
```

Change a value in the matrix

```
mat[2, 5] <- NA
```

Change an entire column

```
mat[, 3] <- 100:103
```

Adding a column or row

```
mat.plus.col <- cbind(mat, 100:103)
mat.plus.row <- rbind(300:307, mat)
```

Warning in rbind(300:307, mat): number of columns of result is not a multiple of vector length (arg 1)

Assign row and column names

```
rownames(mat) <- c("first", "second", "third", "fourth")
colnames(mat) <- letters[1:ncol(mat)]
```

Choose rows and columns by name

```
mat["first", c("e", "c", "d")]
```

```
  e  c  d
17 100 13
```

Choose columns by logical vectors

```
mat[, c(T, T, F, F, T, F)]
```

```
      a b e
first 1 5 17
second 2 6 NA
third 3 7 19
fourth 4 8 20
```

Transpose a matrix

```
t(mat)
```

	first	second	third	fourth
a	1	2	3	4
b	5	6	7	8
c	100	101	102	103
d	13	14	15	16
e	17	NA	19	20
f	21	22	23	24

Add, subtract, multiply, or divide a matrix by a scalar

```
mat * 5
```

	a	b	c	d	e	f
first	5	25	500	65	85	105
second	10	30	505	70	NA	110
third	15	35	510	75	95	115
fourth	20	40	515	80	100	120

```
mat / 3
```

	a	b	c	d	e	f
first	0.3333333	1.666667	33.33333	4.333333	5.666667	7.000000
second	0.6666667	2.000000	33.66667	4.666667	NA	7.333333
third	1.0000000	2.333333	34.00000	5.000000	6.333333	7.666667
fourth	1.3333333	2.666667	34.33333	5.333333	6.666667	8.000000

```
mat ^ 2
```

	a	b	c	d	e	f
first	1	25	10000	169	289	441
second	4	36	10201	196	NA	484
third	9	49	10404	225	361	529
fourth	16	64	10609	256	400	576

Add a column and a matrix

```
mat + 1000:1003
```

	a	b	c	d	e	f
first	1001	1005	1100	1013	1017	1021
second	1003	1007	1102	1015	NA	1023
third	1005	1009	1104	1017	1021	1025
fourth	1007	1011	1106	1019	1023	1027

Row and column sums or means

```
rowSums(mat)
```

	first	second	third	fourth
	157	NA	169	175

```
colMeans(mat)
```

	a	b	c	d	e	f
	2.5	6.5	101.5	14.5	NA	22.5

Arrays

Arrays are multi-dimensional objects that also contain only a single atomic mode of data. They are indexed the same way as matrices, but created by specifying the number of dimensions.

```
# 1 dimensional array (= vector)
```

```
arr.vec <- array(x)
```

```
arr.vec
```

```
[1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
[24] 24
```

```
# 2 dimensional array (= matrix)
```

```
arr.mat <- array(x, dim = c(3, 8))
```

```
arr.mat
```

```
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
[1,]    1    4    7   10   13   16   19   22
[2,]    2    5    8   11   14   17   20   23
[3,]    3    6    9   12   15   18   21   24
```

```
# 3 dimensional array
```

```
arr.3d <- array(x, dim = c(3, 4, 2))
```

```
arr.3d
```

```
, , 1
```

```
      [,1] [,2] [,3] [,4]
[1,]    1    4    7   10
[2,]    2    5    8   11
[3,]    3    6    9   12
```

```
, , 2
```

```
      [,1] [,2] [,3] [,4]
[1,]   13   16   19   22
[2,]   14   17   20   23
[3,]   15   18   21   24
```

Lists

Lists are one-dimensional objects where each dimension can be any kind of object.

```
x <- list(1, letters[1:5], matrix(100:119, 5))
```

```
x
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] "a" "b" "c" "d" "e"
```

```
[[3]]
```

```
      [,1] [,2] [,3] [,4]
[1,]  100  105  110  115
[2,]  101  106  111  116
[3,]  102  107  112  117
```

```
[4,] 103 108 113 118
[5,] 104 109 114 119
```

```
str(x)
```

```
List of 3
```

```
$ : num 1
$ : chr [1:5] "a" "b" "c" "d" ...
$ : int [1:5, 1:4] 100 101 102 103 104 105 106 107 108 109 ...
```

```
class(x)
```

```
[1] "list"
```

```
mode(x)
```

```
[1] "list"
```

A useful piece of information is that lists are special vectors:

```
is.list(x)
```

```
[1] TRUE
```

```
is.vector(x)
```

```
[1] TRUE
```

If you use a single bracket (`[]`) to index a list, you will get a list back:

```
y <- x[2]
str(y)
```

```
List of 1
```

```
$ : chr [1:5] "a" "b" "c" "d" ...
```

```
length(y)
```

```
[1] 1
```

To get the actual object back, you have to use double brackets (`[[]`):

```
z <- x[[2]]
str(z)
```

```
chr [1:5] "a" "b" "c" "d" "e"
```

```
length(z)
```

```
[1] 5
```

List elements can have names and they can be used for indexing like vectors, but single brackets still return a list and double brackets return the object:

```
x2 <- list(first = 1, lets = letters[1:5], third = matrix(30:53, 4))
x2["first"]
```

```
$first
```

```
[1] 1
```

```
x2[["first"]]
```

```
[1] 1
```

The dollar sign (`$`) is a special operator for lists with names that returns the same thing as double brackets:

```
x2$first
```

```
[1] 1
```

List names can be changed with `names`:

```
names(x2) <- c("a.number", "some.letters", "a.matrix")
x2
```

```
$a.number
```

```
[1] 1
```

```
$some.letters
```

```
[1] "a" "b" "c" "d" "e"
```

```
$a.matrix
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]   30   34   38   42   46   50
[2,]   31   35   39   43   47   51
[3,]   32   36   40   44   48   52
[4,]   33   37   41   45   49   53
```

A list can contain a list, and if you know the names, you can chain the `$`:

```
x2$new.element <- list(numbers = 1:5, matrix = matrix(11:25, 3))
x2
```

```
$a.number
```

```
[1] 1
```

```
$some.letters
```

```
[1] "a" "b" "c" "d" "e"
```

```
$a.matrix
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]   30   34   38   42   46   50
[2,]   31   35   39   43   47   51
[3,]   32   36   40   44   48   52
[4,]   33   37   41   45   49   53
```

```
$new.element
```

```
$new.element$numbers
```

```
[1] 1 2 3 4 5
```

```
$new.element$matrix
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]   11   14   17   20   23
[2,]   12   15   18   21   24
[3,]   13   16   19   22   25
```

```
x2$new.element$matrix
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]   11   14   17   20   23
[2,]   12   15   18   21   24
[3,]   13   16   19   22   25
```

To remove an element from a list, you assign `NULL` to that element:

```
x2$some.letters <- NULL
x2
```

```
$a.number
[1] 1
```

```
$a.matrix
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]   30   34   38   42   46   50
[2,]   31   35   39   43   47   51
[3,]   32   36   40   44   48   52
[4,]   33   37   41   45   49   53
```

```
$new.element
$new.element$numbers
[1] 1 2 3 4 5
```

```
$new.element$matrix
      [,1] [,2] [,3] [,4] [,5]
[1,]   11   14   17   20   23
[2,]   12   15   18   21   24
[3,]   13   16   19   22   25
```

Lists can be grown using the `c` function:

```
x <- list(a = 1, b = 2:6, c = letters)
z <- c(x, g = T)
z
```

```
$a
[1] 1
```

```
$b
[1] 2 3 4 5 6
```

```
$c
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q"
[18] "r" "s" "t" "u" "v" "w" "x" "y" "z"
```

```
$g
[1] TRUE
```

We use `dimnames` to add names to arrays. They have to be specified as lists:

```
arr <- array(1:24, dim = c(3, 4, 2))
dimnames(arr) <- list(letters[1:3], LETTERS[1:4], c("one", "two"))
arr
```

```
, , one
```

```
  A B C D
a 1 4 7 10
b 2 5 8 11
c 3 6 9 12
```

```
, , two
```



```

      A  B  C  D
a 13 16 19 22
b 14 17 20 23
c 15 18 21 24

```

Data Frames

Data frames are two-dimensional objects that are normally used to represent data where the rows are observations and the columns are variables.

```

ids <- c(1213, 2435, 5367, 6745, 3592)
loc <- c("north", "north", "north", "west", "south")
len <- c(9.9, 4.5, 7.7, 3.4, 2.0)
wght <- c(270, 130, 235, 90, 88)

df <- data.frame(id = ids, location = loc, len = len, wt = wght)

str(df)

```

```

'data.frame':  5 obs. of  4 variables:
 $ id      : num  1213 2435 5367 6745 3592
 $ location: Factor w/ 3 levels "north","south",...: 1 1 1 3 2
 $ len     : num   9.9  4.5  7.7  3.4  2
 $ wt      : num   270  130  235  90  88

```

```
nrow(df)
```

```
[1] 5
```

```
ncol(df)
```

```
[1] 4
```

Data frames are actually special lists where every column is an element that is the same length:

```
is.data.frame(df)
```

```
[1] TRUE
```

```
is.list(df)
```

```
[1] TRUE
```

```
is.vector(df)
```

```
[1] FALSE
```

```
length(df)
```

```
[1] 4
```

Data frames are indexed the same way as matrices:

```
df[1, ]
```

```

      id location len  wt
1 1213    north 9.9 270

```

```
df[, "len"]
```

```
[1] 9.9 4.5 7.7 3.4 2.0
```

```
df$wt
```

```
[1] 270 130 235 90 88
```

```
df[, c("id", "wt")]
```

```
   id wt
1 1213 270
2 2435 130
3 5367 235
4 6745 90
5 3592 88
```

Data frames are often indexed by a column within the data frame itself. For instance, we want to select only the rows where length is less than 5:

```
df[df$len < 5, ]
```

```
   id location len wt
2 2435    north 4.5 130
4 6745    west 3.4 90
5 3592    south 2.0 88
```

Notice that when we do this, we are placing the condition in the row slot of the indexing brackets. The point of this is that we are creating a logical condition as long as there are rows and using this logical vector to index.

Here's a more complex example:

```
df[df$wt > 200 & df$len < 8, ]
```

```
   id location len wt
3 5367    north 7.7 235
```

We can also choose which columns to return at the same time:

```
df[df$location != "north", c("id", "len", "wt")]
```

```
   id len wt
4 6745 3.4 90
5 3592 2.0 88
```

The `subset` function is a convenient way to index a data.frame without using the `$` notation:

```
subset(df, wt > 200, c("id", "location"))
```

```
   id location
1 1213    north
3 5367    north
```

You can also avoid using the `$` notation by attaching the data frame to the “search path”. This makes the column names in the data frame “findable”:

```
attach(df)
```

The following object is masked `_by_ .GlobalEnv`:

```
len
df[wt > 200, ]
```

```
   id location len wt
```

```
1 1213    north 9.9 270
3 5367    north 7.7 235
```

```
detach(df)
df[location != "north", ]
```

Error in `[.data.frame`(df, location != "north",): object 'location' not found

If you use `attach`, you must remember to use `detach` to remove the data.frame from the search path.

Coercion

Many objects can be coerced from one class to another using `as.<class>` functions. If you have a **numeric** vector, it can be coerced to **character** or **logical**:

```
as.character(1:5)
```

```
[1] "1" "2" "3" "4" "5"
```

when going from numeric to logical, 0 = FALSE, all other numbers are TRUE

```
as.logical(c(-1, -0.5, 0, 1, 3.5, 6))
```

```
[1] TRUE TRUE FALSE TRUE TRUE TRUE
```

Going from character to numeric or logical:

```
as.numeric(c("-5", "0.3", "3.14x", "hello", "a4"))
```

Warning: NAs introduced by coercion

```
[1] -5.0 0.3 NA NA NA
```

```
as.logical(c("hello", "T", "false", "True", "n", "1"))
```

```
[1] NA TRUE FALSE TRUE NA NA
```

Going from logical to character or numeric:

```
as.character(c(T, F, TRUE, FALSE))
```

```
[1] "TRUE" "FALSE" "TRUE" "FALSE"
```

```
as.numeric(c(T, F, TRUE, FALSE))
```

```
[1] 1 0 1 0
```

When coercing a logical to numeric $T = 1$ and $F = 0$. This has some useful properties. To count the number of elements that meet a condition, we can use this feature with the `sum` function:

```
x <- sample(1:5, 100, replace = T)
```

```
x
```

```
[1] 2 1 3 1 1 2 3 5 5 2 3 4 1 1 2 5 2 3 4 4 5 5 1 3 1 1 2 4 1 2 5 4 1 2 5
[36] 5 3 5 5 1 4 5 4 4 1 3 4 5 1 5 5 1 2 5 4 2 3 3 3 1 1 1 5 1 4 2 4 1 1 2
[71] 3 5 5 3 3 3 5 4 5 4 3 5 2 1 4 5 5 3 4 2 5 2 4 3 5 2 5 5 2 5
```

```
sum(x == 1)
```

```
[1] 21
```

Likewise, to calculate the proportion of things that meet a condition, we use the same trick with `mean`:

```
mean(x <= 2)
```

[1] 0.38