

SIOB 296 Introduction to Programming with R

Eric Archer (eric.archer@noaa.gov)

Week 2: April 11, 2017

Reading

The Book of R:

Chapter 3 Matrices and Arrays

Chapter 5 Lists and Data Frames

Chapter 6.2.4 As-Dot Conversion Functions

The Art of R:

Chapter 3 Matrices and Arrays Chapter 4 Lists Chapter 5 Data Frames

sample

The `sample` function is a very useful function for drawing random samples from vectors either with or without replacement.

If it is called with just a vector, it returns a permuted form of that vector:

```
sample(1:10)
```

```
[1] 1 10 2 4 5 7 6 9 3 8
```

A smaller vector can be created by specifying the `size` argument:

```
sample(1:10, size = 5)
```

```
[1] 5 3 10 1 8
```

In this case, 5 unique values are returned. If you want to sample with replacement, specify `replace = TRUE`:

```
sample(1:10, size = 5, replace = TRUE)
```

```
[1] 1 6 7 1 7
```

If a larger vector is to be sampled, then `replace` has to be `TRUE`:

```
sample(1:10, size = 100, replace = TRUE)
```

```
[1] 6 10 5 10 10 1 1 3 1 5 4 6 3 9 2 4 1 10 5 3 7 10 5  
[24] 4 9 10 9 7 2 4 3 7 4 4 2 8 4 2 2 4 4 4 6 7 4 1  
[47] 9 3 3 4 6 8 1 9 7 10 4 5 10 3 10 10 5 10 3 2 3 8 8  
[70] 9 3 2 1 4 9 8 1 5 6 8 8 3 8 7 4 3 1 2 7 8 9 7  
[93] 8 10 3 4 10 4 10 5
```

By default, all elements in the vector have the same likelihood of being sampled. Weights can be applied by specifying them in the `prob` argument:

```
sample(letters[1:5], size = 20, replace = TRUE, prob = c(10, 10, 1, 1, 0))
```

```
[1] "a" "b" "a" "a" "a" "b" "a" "b" "c" "b" "a" "b" "d" "b" "b" "b" "a"  
[18] "a" "a" "b"
```

Matrices

Matrices are always two-dimensional objects having a certain number of rows and columns. They contain only one kind (atomic mode) of data (e.g., numeric, character, logical). They are created by supplying a vector of values to the `matrix()` function and specifying how many rows and/or how many columns to dimension it by

```
# Create a matrix
```

```
x <- 1:24  
mat <- matrix(x, nrow = 4)  
mat
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]  
[1,]    1    5    9   13   17   21  
[2,]    2    6   10   14   18   22  
[3,]    3    7   11   15   19   23  
[4,]    4    8   12   16   20   24
```

```
# How many elements are in the matrix?
```

```
length(mat)
```

```
[1] 24
```

```
#How many rows and columns?
```

```
nrow(mat)
```

```
[1] 4
```

```
ncol(mat)
```

```
[1] 6
```

Cells are selected by [row, column]

```
mat[2, 3]
```

```
[1] 10
```

Selecting a single row or single column returns a vector

```
mat[3, ]
```

```
[1]  3  7 11 15 19 23
```

```
mat[, 4]
```

```
[1] 13 14 15 16
```

Use `drop = F` to select a single row or column and return a matrix

```
mat[4, , drop = F]
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]  
[1,]    4    8   12   16   20   24
```

```
mat[, 2, drop = F]
```

```
      [,1]  
[1,]    5  
[2,]    6  
[3,]    7  
[4,]    8
```

Select several rows or columns

```
mat[c(1, 3, 4), ]
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]     1     5     9    13    17    21
[2,]     3     7    11    15    19    23
[3,]     4     8    12    16    20    24
```

```
mat[, 2:5]
```

```
      [,1] [,2] [,3] [,4]
[1,]     5     9    13    17
[2,]     6    10    14    18
[3,]     7    11    15    19
[4,]     8    12    16    20
```

Select rows, exclude columns

```
mat[1:3, -(2:4)]
```

```
      [,1] [,2] [,3]
[1,]     1    17    21
[2,]     2    18    22
[3,]     3    19    23
```

Change a value in the matrix

```
mat[2, 5] <- NA
```

Change an entire column

```
mat[, 3] <- 100:103
```

Adding a column or row

```
mat.plus.col <- cbind(mat, 100:103)
mat.plus.row <- rbind(300:307, mat)
```

Warning in rbind(300:307, mat): number of columns of result is not a multiple of vector length (arg 1)

Assign row and column names

```
rownames(mat) <- c("first", "second", "third", "fourth")
colnames(mat) <- letters[1:ncol(mat)]
```

Choose rows and columns by name

```
mat["first", c("e", "c", "d")]
```

```
  e   c   d
17 100  13
```

Choose columns by logical vectors

```
mat[, c(T, T, F, F, T, F)]
```

```
      a b e
first 1 5 17
second 2 6 NA
third 3 7 19
fourth 4 8 20
```

Transpose a matrix

```
t(mat)
```

```
  first second third fourth
a      1      2      3      4
b      5      6      7      8
c    100    101    102    103
d     13     14     15     16
e     17     NA     19     20
f     21     22     23     24
```

Add, subtract, multiply, or divide a matrix by a scalar

```
mat * 5
```

```
      a  b  c  d  e  f
first  5 25 500 65  85 105
second 10 30 505 70  NA 110
third  15 35 510 75  95 115
fourth 20 40 515 80 100 120
```

```
mat / 3
```

```
      a      b      c      d      e      f
first 0.3333333 1.666667 33.33333 4.333333 5.666667 7.000000
second 0.6666667 2.000000 33.66667 4.666667      NA 7.333333
third  1.0000000 2.333333 34.00000 5.000000 6.333333 7.666667
fourth 1.3333333 2.666667 34.33333 5.333333 6.666667 8.000000
```

```
mat ^ 2
```

```
      a  b      c  d  e  f
first  1 25 10000 169 289 441
second  4 36 10201 196  NA 484
third   9 49 10404 225 361 529
fourth 16 64 10609 256 400 576
```

Add a column and a matrix

```
mat + 1000:1003
```

```
      a  b      c  d  e  f
first 1001 1005 1100 1013 1017 1021
second 1003 1007 1102 1015  NA 1023
third  1005 1009 1104 1017 1021 1025
fourth 1007 1011 1106 1019 1023 1027
```

Row and column sums or means

```
rowSums(mat)
```

```
  first second  third fourth
    157      NA    169     175
```

```
colMeans(mat)
```

```
      a      b      c      d      e      f
    2.5    6.5  101.5   14.5     NA   22.5
```

Arrays

Arrays are multi-dimensional objects that also contain only a single atomic mode of data. They are indexed the same way as matrices, but created by specifying the number of dimensions.

```
# 1 dimensional array (= vector)
```

```
arr.vec <- array(x)
```

```
arr.vec
```

```
[1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
[24] 24
```

```
# 2 dimensional array (= matrix)
```

```
arr.mat <- array(x, dim = c(3, 8))
```

```
arr.mat
```

```
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
[1,]    1    4    7   10   13   16   19   22
[2,]    2    5    8   11   14   17   20   23
[3,]    3    6    9   12   15   18   21   24
```

```
# 3 dimensional array
```

```
arr.3d <- array(x, dim = c(3, 4, 2))
```

```
arr.3d
```

```
, , 1
```

```
      [,1] [,2] [,3] [,4]
[1,]    1    4    7   10
[2,]    2    5    8   11
[3,]    3    6    9   12
```

```
, , 2
```

```
      [,1] [,2] [,3] [,4]
[1,]   13   16   19   22
[2,]   14   17   20   23
[3,]   15   18   21   24
```

Lists

Lists are one-dimensional objects where each element can be any kind of object.

```
x <- list(1, letters[1:5], matrix(100:119, 5))
```

```
x
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] "a" "b" "c" "d" "e"
```

```
[[3]]
```

```
      [,1] [,2] [,3] [,4]
[1,]  100  105  110  115
[2,]  101  106  111  116
[3,]  102  107  112  117
```

```
[4,] 103 108 113 118
[5,] 104 109 114 119
```

```
str(x)
```

```
List of 3
```

```
$ : num 1
```

```
$ : chr [1:5] "a" "b" "c" "d" ...
```

```
$ : int [1:5, 1:4] 100 101 102 103 104 105 106 107 108 109 ...
```

```
class(x)
```

```
[1] "list"
```

```
mode(x)
```

```
[1] "list"
```

A useful piece of information is that lists are special vectors:

```
is.list(x)
```

```
[1] TRUE
```

```
is.vector(x)
```

```
[1] TRUE
```

If you use a single bracket (`[]`) to index a list, you will get a list back:

```
y <- x[2]
```

```
str(y)
```

```
List of 1
```

```
$ : chr [1:5] "a" "b" "c" "d" ...
```

```
length(y)
```

```
[1] 1
```

To get the actual object back, you have to use double brackets (`[[`):

```
z <- x[[2]]
```

```
str(z)
```

```
chr [1:5] "a" "b" "c" "d" "e"
```

```
length(z)
```

```
[1] 5
```

List elements can have names and they can be used for indexing like vectors, but single brackets still return a list and double brackets return the object:

```
x2 <- list(first = 1, lets = letters[1:5], third = matrix(30:53, 4))
```

```
x2["first"]
```

```
$first
```

```
[1] 1
```

```
x2[["first"]]
```

```
[1] 1
```

The dollar sign (`$`) is a special operator for lists with names that returns the same thing as double brackets:

```
x2$first
```

```
[1] 1
```

List names can be changed with `names`:

```
names(x2) <- c("a.number", "some.letters", "a.matrix")
x2
```

```
$a.number
```

```
[1] 1
```

```
$some.letters
```

```
[1] "a" "b" "c" "d" "e"
```

```
$a.matrix
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]   30   34   38   42   46   50
[2,]   31   35   39   43   47   51
[3,]   32   36   40   44   48   52
[4,]   33   37   41   45   49   53
```

A list can contain a list, and if you know the names, you can chain the `$`:

```
x2$new.element <- list(numbers = 1:5, matrix = matrix(11:25, 3))
x2
```

```
$a.number
```

```
[1] 1
```

```
$some.letters
```

```
[1] "a" "b" "c" "d" "e"
```

```
$a.matrix
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]   30   34   38   42   46   50
[2,]   31   35   39   43   47   51
[3,]   32   36   40   44   48   52
[4,]   33   37   41   45   49   53
```

```
$new.element
```

```
$new.element$numbers
```

```
[1] 1 2 3 4 5
```

```
$new.element$matrix
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]   11   14   17   20   23
[2,]   12   15   18   21   24
[3,]   13   16   19   22   25
```

```
x2$new.element$matrix
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]   11   14   17   20   23
[2,]   12   15   18   21   24
[3,]   13   16   19   22   25
```

To remove an element from a list, you assign `NULL` to that element:

```
x2$some.letters <- NULL
x2
```

```
$a.number
[1] 1
```

```
$a.matrix
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]   30   34   38   42   46   50
[2,]   31   35   39   43   47   51
[3,]   32   36   40   44   48   52
[4,]   33   37   41   45   49   53
```

```
$new.element
$new.element$numbers
[1] 1 2 3 4 5
```

```
$new.element$matrix
      [,1] [,2] [,3] [,4] [,5]
[1,]   11   14   17   20   23
[2,]   12   15   18   21   24
[3,]   13   16   19   22   25
```

Lists can be grown using the `c` function:

```
x <- list(a = 1, b = 2:6, c = letters)
z <- c(x, g = T)
z
```

```
$a
[1] 1
```

```
$b
[1] 2 3 4 5 6
```

```
$c
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q"
[18] "r" "s" "t" "u" "v" "w" "x" "y" "z"
```

```
$g
[1] TRUE
```

We use `dimnames` to add names to arrays. They have to be specified as lists:

```
arr <- array(1:24, dim = c(3, 4, 2))
dimnames(arr) <- list(letters[1:3], LETTERS[1:4], c("one", "two"))
arr
```

```
, , one
```

```
  A B C D
a 1 4 7 10
b 2 5 8 11
c 3 6 9 12
```

```
, , two
```



```

      A  B  C  D
a 13 16 19 22
b 14 17 20 23
c 15 18 21 24

```

Data Frames

Data frames are two-dimensional objects that are normally used to represent data where the rows are observations and the columns are variables.

```

ids <- c(1213, 2435, 5367, 6745, 3592)
loc <- c("north", "north", "north", "west", "south")
len <- c(9.9, 4.5, 7.7, 3.4, 2.0)
wght <- c(270, 130, 235, 90, 88)

df <- data.frame(id = ids, location = loc, len = len, wt = wght)

str(df)

```

```

'data.frame':  5 obs. of  4 variables:
 $ id      : num  1213 2435 5367 6745 3592
 $ location: Factor w/ 3 levels "north","south",...: 1 1 1 3 2
 $ len     : num   9.9  4.5  7.7  3.4  2
 $ wt      : num   270  130  235  90  88

```

```
nrow(df)
```

```
[1] 5
```

```
ncol(df)
```

```
[1] 4
```

Data frames are actually special lists where every column is an element that is the same length:

```
is.data.frame(df)
```

```
[1] TRUE
```

```
is.list(df)
```

```
[1] TRUE
```

```
is.vector(df)
```

```
[1] FALSE
```

```
length(df)
```

```
[1] 4
```

Data frames are indexed the same way as matrices:

```
df[1, ]
```

```

      id location len  wt
1 1213    north 9.9 270

```

```
df[, "len"]
```

```
[1] 9.9 4.5 7.7 3.4 2.0
```

```
df[, c("id", "wt")]
```

```
   id wt
1 1213 270
2 2435 130
3 5367 235
4 6745  90
5 3592  88
```

Columns can also be returned as a vector using the \$:

```
df$wt
```

```
[1] 270 130 235  90  88
```

Data frames are often indexed by a column within the data frame itself. For instance, we want to select only the rows where length is less than 5:

```
df[df$len < 5, ]
```

```
   id location len wt
2 2435   north 4.5 130
4 6745   west 3.4  90
5 3592   south 2.0  88
```

Notice that when we do this, we are placing the condition in the row slot of the indexing brackets. The point of this is that we are creating a logical condition as long as there are rows and using this logical vector to index.

Here's a more complex example:

```
df[df$wt > 200 & df$len < 8, ]
```

```
   id location len wt
3 5367   north 7.7 235
```

We can also choose which columns to return at the same time:

```
df[df$location != "north", c("id", "len", "wt")]
```

```
   id len wt
4 6745 3.4 90
5 3592 2.0 88
```

The `subset` function is a convenient way to index a data.frame without using the \$ notation:

```
subset(df, wt > 200, c("id", "location"))
```

```
   id location
1 1213   north
3 5367   north
```

You can also avoid using the \$ notation by **attaching** the data frame to the “search path”. This makes the column names in the data frame “findable”:

```
attach(df)
```

The following object is masked `_by_ .GlobalEnv`:

```
len
df[wt > 200, ]
```

```

      id location len wt
1 1213    north 9.9 270
3 5367    north 7.7 235

```

```

detach(df)
df[location != "north", ]

```

Error in `[.data.frame`(df, location != "north",) : object 'location' not found

If you use `attach`, you must remember to use `detach` to remove the data.frame from the search path. **Avoid using `attach`!!! It opens the potential for sneaky errors.**

Coercion

Many objects can be coerced from one class to another using `as.<class>` functions. If you have a numeric vector, it can be coerced to character or logical:

```
as.character(1:5)
```

```
[1] "1" "2" "3" "4" "5"
```

when going from numeric to logical, 0 = FALSE, all other numbers are TRUE

```
as.logical(c(-1, -0.5, 0, 1, 3.5, 6))
```

```
[1] TRUE TRUE FALSE TRUE TRUE TRUE
```

Going from character to numeric or logical:

```
as.numeric(c("-5", "0.3", "3.14x", "hello", "a4"))
```

Warning: NAs introduced by coercion

```
[1] -5.0 0.3 NA NA NA
```

```
as.logical(c("hello", "T", "false", "True", "n", "1"))
```

```
[1] NA TRUE FALSE TRUE NA NA
```

Going from logical to character or numeric:

```
as.character(c(T, F, TRUE, FALSE))
```

```
[1] "TRUE" "FALSE" "TRUE" "FALSE"
```

```
as.numeric(c(T, F, TRUE, FALSE))
```

```
[1] 1 0 1 0
```

When coercing a logical to numeric $T = 1$ and $F = 0$. This has some useful properties. To count the number of elements that meet a condition, we can use this feature with the `sum` function:

```
x <- sample(1:5, 100, replace = T)
```

```
x
```

```

[1] 4 5 4 4 4 4 1 2 4 1 3 1 2 1 4 1 4 1 2 5 2 5 4 2 3 1 1 2 2 3 3 4 1 5 1
[36] 5 2 5 5 5 2 1 4 3 4 4 4 2 2 1 1 3 4 5 1 4 2 2 3 3 3 4 4 2 4 5 5 3 5 1
[71] 2 4 5 2 3 4 3 2 5 5 1 4 1 5 5 4 5 5 4 4 2 4 3 1 5 4 5 1 5 4

```

```
sum(x == 1)
```

```
[1] 19
```

Likewise, to calculate the proportion of things that meet a condition, we use the same trick with `mean`:

```
mean(x <= 2)
```

```
[1] 0.37
```