

# SIOB 296 Introduction to Programming with R

*Eric Archer (eric.archer@noaa.gov)*

*Week 4: April 25, 2017*

---

## Reading

*The Book of R:*

Chapter 4.2: Characters

*The Art of R:*

Chapter 11.1, 11.2: String Manipulation

Chapter 3.3: Applying Functions to Matrix Rows and Columns

Chapter 6.2: Common Functions Used with Factors

## Merging data frames

To merge two data frames based on common columns, use the `merge` function. To demonstrate how it works, we'll load several simple data frames that have data about a related set of specimens:

```
spec.dfs <- load("merge_data.rdata")
print(spec.dfs)
```

```
[1] "specimens" "cranial"    "haps"      "sex"       "trials"
```

```
# Skull measurements
```

```
str(cranial)
```

```
'data.frame':  20 obs. of  2 variables:
```

```
 $ id   : chr  "Specimen-25" "Specimen-37" "Specimen-4" "Specimen-1" ...
```

```
 $ skull: num  257 251 261 256 259 ...
```

```
# Haplotypes
```

```
str(haps)
```

```
'data.frame':  20 obs. of  2 variables:
```

```
 $ id   : chr  "Specimen-7" "Specimen-31" "Specimen-1" "Specimen-32" ...
```

```
 $ haps: chr   "E" "C" "D" "E" ...
```

```
# Sex
```

```
str(sex)
```

```
'data.frame':  40 obs. of  2 variables:
```

```
 $ specimens: chr  "Specimen-1" "Specimen-2" "Specimen-3" "Specimen-4" ...
```

```
 $ sex      : chr  "M" "F" "M" "F" ...
```

```
# Results of multiple trials
```

```
str(trials)
```

```
'data.frame':  30 obs. of  2 variables:
```

```
 $ id   : chr  "Specimen-27" "Specimen-32" "Specimen-13" "Specimen-29" ...
```

```
 $ value: num  20.3 23.6 20.3 22.2 38.1 ...
```

`merge` requires two data frames as its first two arguments, `x` and `y`. The two primary sets of arguments that control how the merging happens are `by` (with the related `by.x` and `by.y`) which identifies which column(s) are used to identify equivalent records to be merged, and `all` (with the related `all.x` and `all.y`) which specifies whether the data frame returned contains all records from both data.frames or just all records from one or the other. As an example, we'll merge the skull measurement and haplotype data frames. They have the same number of rows, but not the exact same specimens. Because they both have a column in common (`id`), by default, it is used for by:

```
merge(cranial, haps)
```

	id	skull	haps
1	Specimen-1	255.7784	D
2	Specimen-10	260.8000	B
3	Specimen-15	262.2992	E
4	Specimen-21	255.7237	A
5	Specimen-25	257.3424	D
6	Specimen-29	271.5665	C
7	Specimen-32	263.1267	E
8	Specimen-37	251.4024	E
9	Specimen-4	261.2716	A
10	Specimen-6	264.0681	B

Note that there are only 10 rows because by default, the `all` argument is set to `FALSE` which means the function will only return a data frame with ids that occur in both original data frames. If we want to return a data frame with all specimens in both original data frames, we set `all = TRUE`:

```
merge(cranial, haps, all = TRUE)
```

	id	skull	haps
1	Specimen-1	255.7784	D
2	Specimen-10	260.8000	B
3	Specimen-11	NA	E
4	Specimen-12	261.7318	<NA>
5	Specimen-13	NA	A
6	Specimen-15	262.2992	E
7	Specimen-17	264.1644	<NA>
8	Specimen-18	NA	B
9	Specimen-19	255.9717	<NA>
10	Specimen-21	255.7237	A
11	Specimen-22	NA	B
12	Specimen-23	270.7938	<NA>
13	Specimen-24	264.6283	<NA>
14	Specimen-25	257.3424	D
15	Specimen-26	260.7033	<NA>
16	Specimen-28	264.2351	<NA>
17	Specimen-29	271.5665	C
18	Specimen-31	NA	C
19	Specimen-32	263.1267	E
20	Specimen-33	261.7152	<NA>
21	Specimen-34	253.3402	<NA>
22	Specimen-35	NA	C
23	Specimen-36	258.5202	<NA>
24	Specimen-37	251.4024	E
25	Specimen-38	NA	B
26	Specimen-4	261.2716	A
27	Specimen-5	NA	D

```

28 Specimen-6 264.0681    B
29 Specimen-7      NA    E
30 Specimen-9      NA    D

```

Note that here, NAs are inserted where there is no data in the other data frame. We can also specify that we only want all records in one data frame:

```
merge(cranial, haps, all.x = TRUE)
```

```

      id    skull haps
1 Specimen-1 255.7784    D
2 Specimen-10 260.8000    B
3 Specimen-12 261.7318 <NA>
4 Specimen-15 262.2992    E
5 Specimen-17 264.1644 <NA>
6 Specimen-19 255.9717 <NA>
7 Specimen-21 255.7237    A
8 Specimen-23 270.7938 <NA>
9 Specimen-24 264.6283 <NA>
10 Specimen-25 257.3424    D
11 Specimen-26 260.7033 <NA>
12 Specimen-28 264.2351 <NA>
13 Specimen-29 271.5665    C
14 Specimen-32 263.1267    E
15 Specimen-33 261.7152 <NA>
16 Specimen-34 253.3402 <NA>
17 Specimen-36 258.5202 <NA>
18 Specimen-37 251.4024    E
19 Specimen-4 261.2716    A
20 Specimen-6 264.0681    B

```

If the common column is not the same in both data frames, you have to specify it with `by.x` and `by.y`:

```
merged.df <- merge(sex, trials, by.x = "specimens", by.y = "id")
head(merged.df)
```

```

      specimens sex    value
1 Specimen-11    M 28.29490
2 Specimen-12    M 36.97172
3 Specimen-12    M 25.35317
4 Specimen-13    M 26.55486
5 Specimen-13    M 20.30126
6 Specimen-14    F 31.92108

```

Here, the identifier column name of the `x` data.frame is retained.

## Character and string manipulation

### nchar

A character vector is a vector where every element is a character string of any length. The `length()` of a character vector is the number of elements in it:

```
x <- c("This is a sentence", "Hello World!", "This is the third element")
length(x)
```

```
[1] 3
```

To get the number of characters in each element, use `nchar()`:

```
nchar(x)
```

```
[1] 18 12 25
```

## substr

Strings can be extracted from elements using `substr()`. You specify the first and last characters to be extracted from each string:

```
# get the first three characters from every string  
substr(x, 1, 3)
```

```
[1] "Thi" "Hel" "Thi"
```

```
# get the 3rd character from every string  
substr(x, 3, 3)
```

```
[1] "i" "l" "i"
```

`substr` can also be used to replace values within strings by assigning:

```
substr(x, 1, 4) <- "That"  
x
```

```
[1] "That is a sentence"      "Thato World!"  
[3] "That is the third element"
```

## strsplit

Strings can be split based on some common delimiter using `strsplit()`:

```
# split based on spaces  
x.split <- strsplit(x, " ")  
x.split
```

```
[[1]]  
[1] "That"      "is"        "a"         "sentence"  
  
[[2]]  
[1] "Thato"    "World!"  
  
[[3]]  
[1] "That"      "is"        "the"       "third"     "element"
```

```
str(x.split)
```

```
List of 3  
 $ : chr [1:4] "That" "is" "a" "sentence"  
 $ : chr [1:2] "Thato" "World!"  
 $ : chr [1:5] "That" "is" "the" "third" ...
```

Note that the return value from `strsplit` is a list. Each element in the list corresponds to a vector resulting from splitting every element in the original vector

```
x.split[[1]]
```

```
[1] "That"      "is"        "a"         "sentence"
```

## paste

To create strings from combinations of strings (or numbers) we use `paste()`. This function takes a set of vectors, and pastes the elements together using recycling:

```
# vectors are equal length  
paste(letters[1:6], 1:6)
```

```
[1] "a 1" "b 2" "c 3" "d 4" "e 5" "f 6"
```

```
# one vector is a multiple of the other  
paste(letters[1:6], 1:2)
```

```
[1] "a 1" "b 2" "c 1" "d 2" "e 1" "f 2"
```

```
# one vector is not a multiple of the other  
paste(letters[1:6], 1:4)
```

```
[1] "a 1" "b 2" "c 3" "d 4" "e 1" "f 2"
```

The argument `sep` determines what character is used as a separator between the characters:

```
paste(letters[1:6], 1:2, sep = "-")
```

```
[1] "a-1" "b-2" "c-1" "d-2" "e-1" "f-2"
```

If you do not want a separator character, either set `sep = ""` or use `paste0()`:

```
paste0(letters[1:6], 1:2)
```

```
[1] "a1" "b2" "c1" "d2" "e1" "f2"
```

If you want to paste all of the arguments to create a single element vector, set the `collapse` argument:

```
paste(letters[1:6], 1:2, sep = "-", collapse = "#")
```

```
[1] "a-1#b-2#c-1#d-2#e-1#f-2"
```

## tolower, toupper

Character case can be changed with `tolower` and `toupper`:

```
tolower(x)
```

```
[1] "that is a sentence"      "thato world!"  
[3] "that is the third element"
```

```
toupper(x)
```

```
[1] "THAT IS A SENTENCE"      "THATO WORLD!"  
[3] "THAT IS THE THIRD ELEMENT"
```

## Regular Expressions

For finer control on searching and replacing text within strings, you will have to turn to “regular expressions”, which is a kind of syntax of its own and is common across several platforms. The help page for regular expressions in R is `?regex`. The functions that are most commonly used with regular expressions are given in `grep`. The most commonly used on this page are:

`grep` and `grep1`: Identify elements that have the sought after pattern `sub` and `gsub`: Replace a desired pattern with other text

```
x <- c("Here is some text", "This is more text", "I have the number 1", "22 is the number I have")
# which elements have the word "text"?
grep("text", x)
```

```
[1] 1 2
```

```
# which elements have numbers?
grep("[[:digit:]]", x)
```

```
[1] 3 4
```

```
# replace the word "This" with "That"
gsub("This", "That", x)
```

```
[1] "Here is some text"      "That is more text"
[3] "I have the number 1"    "22 is the number I have"
```

## apply Functions

Many times, we want to execute the same function on sequential elements of some object. This could be things like the elements of a vector or list, the rows of a matrix, or the columns of a data frame. For these, R provides a family of functions that usually end in `-apply` or are based on them.

### `lapply`

The most basic of these functions is `lapply`. The “l” refers to the fact that `lapply` will always return a list. There are two main arguments to `lapply`: the first is the object to be iterated over, and the second is a function that takes sequential elements of that object. As an example, let’s use the `sample` function. Recall that if you execute `sample` with a single integer(`n`), it will return a permutation of the vector `1:n`:

```
sample(5)
```

```
[1] 2 4 5 3 1
```

```
sample(10)
```

```
[1] 4 2 3 1 5 8 9 7 10 6
```

Here is a list resulting from calls to `sample` with the elements of the vector `1:5`:

```
x <- lapply(1:5, sample)
str(x)
```

```
List of 5
```

```
$ : int 1
$ : int [1:2] 2 1
$ : int [1:3] 2 1 3
$ : int [1:4] 3 1 4 2
$ : int [1:5] 5 3 2 4 1
```

```
x
```

```
[[1]]
[1] 1
```

```
[[2]]
[1] 2 1
```

```
[[3]]  
[1] 2 1 3
```

```
[[4]]  
[1] 3 1 4 2
```

```
[[5]]  
[1] 5 3 2 4 1
```

Note that the result is a list, the elements of which are the result of calls to `sample(1)`, `sample(2)`, `sample(3)`, etc. The elements of the return value are in the same order as the original object being iterated over:

```
lapply(c(5, 3, 1, 8), sample)
```

```
[[1]]  
[1] 3 1 5 2 4
```

```
[[2]]  
[1] 3 1 2
```

```
[[3]]  
[1] 1
```

```
[[4]]  
[1] 8 5 6 4 3 2 7 1
```

The first argument can be a list too:

```
lapply(x, sum)
```

```
[[1]]  
[1] 1
```

```
[[2]]  
[1] 3
```

```
[[3]]  
[1] 6
```

```
[[4]]  
[1] 10
```

```
[[5]]  
[1] 15
```

## sapply

If the return value from every iteration was the same length, you may want to simplify the result. This is what `sapply` is for. If every call to the function returns a scalar, then `sapply` will return a vector. If every call to the function returns a vector of equal length, then `sapply` will return a matrix. If every call to the function returns a value of different lengths, then `sapply` defaults to returning a list:

```
# every return value from sum is a scalar - sapply returns a vector  
sapply(x, sum)
```

```
[1] 1 3 6 10 15
```

```
# every return value from sample is a 5 element vector - sapply returns a matrix
sapply(rep(5, 8), sample)
```

```
      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8]
[1,]    4    1    1    4    2    4    2    1
[2,]    5    5    3    3    3    3    3    5
[3,]    1    2    4    1    4    1    1    2
[4,]    3    3    5    2    1    5    5    3
[5,]    2    4    2    5    5    2    4    4
```

```
# this is the same as our lapply example - sapply returns a list
sapply(c(5, 3, 1, 8), sample)
```

```
[[1]]
[1] 1 4 3 5 2
```

```
[[2]]
[1] 1 2 3
```

```
[[3]]
[1] 1
```

```
[[4]]
[1] 1 4 2 7 6 3 8 5
```

Arguments to the function can be specified in the `lapply` or `sapply` call:

```
sapply(c(5, 3, 1, 8), sample, size = 5, replace = TRUE)
```

```
      [,1] [,2] [,3] [,4]
[1,]    4    1    1    7
[2,]    4    3    1    2
[3,]    2    2    1    7
[4,]    2    2    1    6
[5,]    3    1    1    7
```

## apply

If you are dealing with a multi-dimensional object (matrix, array, or data frame) and you want to apply a function to a given dimension (i.e, each row or each column), use `apply`. You have to specify the dimension that you will be iterating over as the second argument (1 = rows, 2 = columns, etc). `apply` will try to simplify the results like `sapply`:

```
x <- matrix(sample(1:100, 24, replace = TRUE), nrow = 4)
x
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]   61   91    6   31   18   38
[2,]   46   81   56   50   45   83
[3,]    3   75   40   51   79    8
[4,]   76   40   94   79   22   47
```

```
# median of each row
apply(x, 1, median)
```

```
[1] 34.5 53.0 45.5 61.5
```



```
# difference of each column
apply(x, 2, diff)
```

```
      [,1] [,2] [,3] [,4] [,5] [,6]
[1,]  -15  -10   50   19   27   45
[2,]  -43   -6  -16    1   34  -75
[3,]   73  -35   54   28  -57   39
```

## tapply

If you want to execute a function on groups of values, **tapply** can often be a good choice. The arguments are a vector that will be summarized, another vector or set of vectors that represent identify elements to groups, and the function that will get the sequential subsets of the original vector. As a simple example, we calculate the means of subsets of a random vector of numbers

```
x <- sample(1:100, 100, replace = TRUE)
grp <- sample(letters[1:5], 100, replace = TRUE)
tapply(x, grp, mean)
```

```
      a      b      c      d      e
44.80000 44.50000 54.61905 66.22222 48.15000
```

As a more practical example, we can calculate the average temperature at each station in our ctd dataset:

```
ctd <- read.csv("ctd.csv")
tapply(ctd$temp, ctd$station, mean)
```

```
Station.1 Station.10 Station.11 Station.12 Station.13 Station.14
 13.56772  14.57675   14.93466   14.40093   13.84903   14.42424
Station.15 Station.16 Station.17 Station.18 Station.19 Station.2
 14.10233   14.42280   14.27389   14.57620   16.32791   14.24618
Station.20 Station.21 Station.22 Station.23 Station.24 Station.25
 13.34375   13.77648   14.18930   14.60143   16.35819   15.82101
Station.26 Station.27 Station.28 Station.29 Station.3 Station.30
 15.84702   14.10141   13.21872   13.76510   14.23830   14.19600
Station.31 Station.32 Station.33 Station.34 Station.35 Station.36
 14.85625   16.73097   14.07543   14.62647   14.74482   15.53890
Station.37 Station.38 Station.39 Station.4 Station.40 Station.5
 15.12451   15.57174   15.14841   14.51093   16.37120   14.65009
Station.6 Station.7 Station.8 Station.9
 14.23170   13.63123   13.94914   14.18727
```

We can use two grouping variables to return a matrix. However, when we do this, the second argument must be specified as a list.

```
# What is the average temperature at each station and depth?
mean.temp <- tapply(ctd$temp, list(station = ctd$station, depth = ctd$depth), mean)
head(mean.temp[, 1:5])
```

```
      depth
station 1      2      3      4      5
Station.1 17.22627 17.18102 17.07373 16.92864 16.75797
Station.10 16.67695 16.55271 16.24712 15.95271 15.66983
Station.11 16.39310 16.17458 15.89458 15.59407 15.34814
Station.12 16.86448 16.74119 16.52642 16.35448 16.15030
Station.13 17.05638 16.86203 16.68034 16.49610 16.24271
```

```
Station.14 16.98061 16.81493 16.59866 16.41940 16.16791
```

## aggregate

If we want to apply the same grouped summary to every column in a data frame, we can use `aggregate`:

```
# what is the median of each measurement at each station?
st.medians <- aggregate(ctd[, 3:8], list(station = ctd$station), median, na.rm = TRUE)
head(st.medians)
```

	station	temp	salinity	dox	ph	pct_light	density
1	Station.1	13.070	33.4570	7.05	8.05	88.330	25.1380
2	Station.10	14.445	33.4695	7.90	8.15	81.530	24.8515
3	Station.11	14.940	33.4625	7.88	8.15	76.270	24.7710
4	Station.12	14.095	33.4530	7.66	8.12	85.035	24.8905
5	Station.13	13.500	33.4635	7.42	8.10	86.640	25.0325
6	Station.14	14.170	33.4640	7.67	8.13	84.700	24.8860

Be careful if the function returns more than one thing though.

```
st.range <- aggregate(ctd[, 3:8], list(station = ctd$station), range, na.rm = TRUE)
head(st.range)
```

	station	temp.1	temp.2	salinity.1	salinity.2	dox.1	dox.2	ph.1	ph.2
1	Station.1	9.92	22.74	33.130	34.033	2.06	10.61	7.66	8.62
2	Station.10	10.36	22.65	33.162	33.864	2.14	13.03	7.66	8.55
3	Station.11	10.58	23.06	33.209	33.817	2.52	11.77	7.69	8.50
4	Station.12	10.24	23.00	32.561	34.311	2.28	11.38	7.67	8.63
5	Station.13	10.00	22.99	33.090	33.879	2.51	10.88	7.69	8.59
6	Station.14	10.20	22.74	33.069	33.891	2.25	11.35	7.65	8.61

  

	pct_light.1	pct_light.2	density.1	density.2
1	69.45	92.25	22.923	26.196
2	30.53	89.64	22.945	25.995
3	5.34	89.20	22.822	25.895
4	47.59	90.87	22.841	26.041
5	55.29	91.79	22.841	26.076
6	41.77	90.71	22.909	26.052

```
str(st.range)
```

```
'data.frame': 40 obs. of 7 variables:
 $ station : Factor w/ 40 levels "Station.1","Station.10",...: 1 2 3 4 5 6 7 8 9 10 ...
 $ temp    : num [1:40, 1:2] 9.92 10.36 10.58 10.24 10 ...
 $ salinity: num [1:40, 1:2] 33.1 33.2 33.2 32.6 33.1 ...
 $ dox     : num [1:40, 1:2] 2.06 2.14 2.52 2.28 2.51 2.25 2.26 2.24 2.33 2.33 ...
 $ ph      : num [1:40, 1:2] 7.66 7.66 7.69 7.67 7.69 7.65 7.67 7.66 7.65 7.67 ...
 $ pct_light: num [1:40, 1:2] 69.45 30.53 5.34 47.59 55.29 ...
 $ density : num [1:40, 1:2] 22.9 22.9 22.8 22.8 22.8 ...
```

Note that the column names seem to have `.1` and `.2` when you print the object, but they aren't in the structure. In this case, every measurement column is itself a two column matrix:

```
dim(st.range$temp)
```

```
[1] 40 2
```

```
head(st.range$temp)
```

```

      [,1] [,2]
[1,]  9.92 22.74
[2,] 10.36 22.65
[3,] 10.58 23.06
[4,] 10.24 23.00
[5,] 10.00 22.99
[6,] 10.20 22.74

```

## by

To apply a function to an entire data frame, use `by()`, which works much like `tapply()`:

```
# How many records per station?
```

```
st.rows <- by(ctd, ctd$station, nrow)
head(st.rows)
```

```
ctd$station
Station.1 Station.10 Station.11 Station.12 Station.13 Station.14
      3535      1120       762      1876      2229      1865
```

```
str(st.rows)
```

```

'by' int [1:40(1d)] 3535 1120 762 1876 2229 1865 1826 1865 1474 1120 ...
- attr(*, "dimnames")=List of 1
  ..$ ctd$station: chr [1:40] "Station.1" "Station.10" "Station.11" "Station.12" ...
- attr(*, "call")= language by.data.frame(data = ctd, INDICES = ctd$station, FUN = nrow)

```

```
st.rows["Station.5"]
```

```

Station.5
      809

```

You can also summarize with multiple groups, which have to be included as a list:

```
# How many records per station?
```

```
st.depth.rows <- by(ctd, list(station = ctd$station, depth = ctd$depth), nrow)
str(st.depth.rows)
```

```

'by' int [1:40, 1:60] 59 59 58 67 58 66 59 67 59 59 ...
- attr(*, "dimnames")=List of 2
  ..$ station: chr [1:40] "Station.1" "Station.10" "Station.11" "Station.12" ...
  ..$ depth : chr [1:60] "1" "2" "3" "4" ...
- attr(*, "call")= language by.data.frame(data = ctd, INDICES = list(station = ctd$station, depth = ctd$depth), FUN = nrow)

```

```
# The object can be indexed like a matrix
```

```
st.depth.rows["Station.1", "12"]
```

```
[1] 59
```

## mapply

To apply a function to sequential elements of multiple vectors, use `mapply()`. The first argument is a function, and every argument afterwards is an argument to that function composed of vectors being iterated over. For example, the following creates a list of random numbers of alternating length with increasing range:

```
mapply(sample, x = 5:10, size = c(20, 4), replace = TRUE)
```

```
[[1]]
[1] 3 1 4 2 4 5 5 4 5 4 2 5 2 3 2 4 4 2 5 3
```

```
[[2]]
[1] 3 5 4 6
```

```
[[3]]
[1] 3 7 3 6 5 3 2 2 1 7 2 1 5 6 1 3 6 5 2 6
```

```
[[4]]
[1] 8 4 6 4
```

```
[[5]]
[1] 9 6 3 2 1 1 7 3 5 4 9 4 8 7 8 7 9 6 8 8
```

```
[[6]]
[1] 5 3 6 4
```

## split

A handy function for creating lists based on a grouping variable is `split`. It will split a vector, matrix, or data frame. For instance, here is a list where every element is a data frame containing only one station's data:

```
st.list <- split(ctd, ctd$station)
head(st.list[[1]])
```

	station	sample_date	temp	salinity	dox	ph	pct_light	density	depth
1	Station.1	2012-11-08	16.81	33.420	8.07	8.20	90.32	24.346	16
2	Station.1	2012-04-19	10.52	33.805	3.16	7.73	88.14	25.930	18
3	Station.1	2010-01-06	15.11	33.415	7.22	8.13	88.97	24.725	32
4	Station.1	2014-02-06	14.00	33.430	7.31	NA	88.01	24.974	41
5	Station.1	2011-01-05	14.20	33.286	7.91	8.16	86.17	24.822	3
6	Station.1	2015-02-03	13.92	33.382	6.45	8.05	87.68	24.953	51

```
head(st.list[[2]])
```

	station	sample_date	temp	salinity	dox	ph	pct_light	density
3536	Station.10	2010-05-10	14.99	33.479	9.62	8.35	70.32	24.799
3537	Station.10	2011-02-02	13.10	33.337	7.24	8.06	65.39	25.085
3538	Station.10	2010-03-17	13.45	33.406	8.62	8.17	73.64	25.069
3539	Station.10	2016-08-02	19.91	33.465	8.98	8.28	82.14	23.616
3540	Station.10	2016-11-02	14.00	33.279	6.68	8.02	79.46	24.858
3541	Station.10	2010-03-17	13.53	33.404	8.62	8.20	72.08	25.050

depth

3536	4
3537	6
3538	6
3539	12
3540	19
3541	4

Here's the same creating an element for each cast (station x date):

```
st.dt.list <- split(ctd, list(station = ctd$station, date = ctd$sample_date))
st.dt.list[[1]]
```

```

[1] station      sample_date temp      salinity  dox      ph
[7] pct_light    density    depth
<0 rows> (or 0-length row.names)

```

Because it does all combinations of the grouping factors, a lot will be empty. Let's find them:

```

num.rows <- sapply(st.dt.list, nrow)
zero.rows <- which(num.rows == 0)
st.dt.list <- st.dt.list[-zero.rows]
st.dt.list[[1]]

```

	station	sample_date	temp	salinity	dox	ph	pct_light	density
5929	Station.12	2010-01-05	14.72	33.374	7.61	8.18	78.87	24.779
6294	Station.12	2010-01-05	14.72	33.374	7.61	8.18	79.14	24.778
6295	Station.12	2010-01-05	14.72	33.373	7.59	8.18	79.32	24.778
6750	Station.12	2010-01-05	14.72	33.373	7.59	8.18	79.11	24.778
6775	Station.12	2010-01-05	14.72	33.374	7.60	8.18	78.97	24.779
6778	Station.12	2010-01-05	14.64	33.375	7.50	8.18	77.96	24.796
6794	Station.12	2010-01-05	14.72	33.373	7.59	8.18	79.13	24.778
6856	Station.12	2010-01-05	14.71	33.374	7.56	8.18	79.24	24.781
6957	Station.12	2010-01-05	14.59	33.373	7.43	8.16	68.05	24.807
6973	Station.12	2010-01-05	14.59	33.373	7.44	8.16	69.51	24.806
6992	Station.12	2010-01-05	14.72	33.365	7.59	8.18	78.82	24.773
7061	Station.12	2010-01-05	14.71	33.363	7.62	8.19	78.36	24.772
7067	Station.12	2010-01-05	14.75	33.245	7.52	8.17	74.31	24.674
7087	Station.12	2010-01-05	14.72	33.368	7.58	8.18	79.09	24.775
7094	Station.12	2010-01-05	14.59	33.373	7.45	8.17	70.50	24.805
7103	Station.12	2010-01-05	14.60	33.374	7.45	8.17	72.98	24.804
7108	Station.12	2010-01-05	14.75	33.243	7.52	8.17	74.02	24.671
7131	Station.12	2010-01-05	14.73	33.328	7.61	8.18	76.89	24.742
7161	Station.12	2010-01-05	14.71	33.364	7.61	8.18	79.00	24.773
7172	Station.12	2010-01-05	14.71	33.360	7.64	8.18	78.11	24.770
7199	Station.12	2010-01-05	14.72	33.366	7.58	8.18	78.80	24.773
7218	Station.12	2010-01-05	14.72	33.347	7.63	8.18	77.19	24.759
7260	Station.12	2010-01-05	14.73	33.310	7.58	8.18	74.98	24.727
7266	Station.12	2010-01-05	14.72	33.372	7.59	8.18	78.90	24.777
7271	Station.12	2010-01-05	14.72	33.368	7.60	8.18	78.87	24.775
7275	Station.12	2010-01-05	14.74	33.268	7.55	8.17	73.94	24.693
7284	Station.12	2010-01-05	14.73	33.324	7.60	8.18	76.35	24.739
7285	Station.12	2010-01-05	14.72	33.371	7.58	8.18	79.13	24.777

	depth
5929	21
6294	20
6295	19
6750	18
6775	22
6778	24
6794	17
6856	23
6957	28
6973	27
6992	11
7061	9
7067	2
7087	14

7094	26
7103	25
7108	1
7131	6
7161	10
7172	8
7199	12
7218	7
7260	4
7266	16
7271	13
7275	3
7284	5
7285	15