

SIOB 296 Introduction to Programming with R

Eric Archer (eric.archer@noaa.gov)

Week 7 (May 16, 2017) - Graphics

Reading

The Book of R:

Chapter 18 - Hypothesis testing

Chapter 20 - Simple linear regression

The Art of R:

Chapter 8.2 - Functions for statistical distributions

Chapter 8.6 - Simulation programming in R

Distributions

Functions are provided to calculate the density, distribution function, quantile function, and generate random numbers from a variety of parametric distributions. They have similar forms, where if `<stat>` is the name of the distribution (e.g., `norm` for Normal, `unif` for Uniform, `binom` for Binomial), `d<stat>` gives the density or probability mass function (likelihood), `p<stat>` gives the probability distribution (cumulative distribution function), `q<stat>` gives the quantile function, and `r<stat>` generates random numbers. Here are examples of all for for the Normal distribution:

```
# The likelihood of five values in a Normal distribution with a
# mean of 10 and a standard deviation of 2:
x <- c(5, 8, 10, 12, 15)
dnorm(x, mean = 10, sd = 2)
```

```
[1] 0.00876415 0.12098536 0.19947114 0.12098536 0.00876415
```

```
# Cumulative probability of same values:
pnorm(x, mean = 10, sd = 2)
```

```
[1] 0.006209665 0.158655254 0.500000000 0.841344746 0.993790335
```

```
# Quantiles:
p <- c(0.05, 0.25, 0.5, 0.75, 0.95)
qnorm(p, mean = 10, sd = 2)
```

```
[1] 6.710293 8.651020 10.000000 11.348980 13.289707
```

```
# Five random draws:
rnorm(5, mean = 10, sd = 2)
```

```
[1] 9.900195 10.892659 10.708232 12.521264 10.287962
```

The random number seed is set with `set.seed()`. Setting this value ensures that the same random number sequence will be repeated:

```
# repeat the same random 5 numbers
set.seed(1)
rnorm(5, mean = 10, sd = 2)
```

```
[1] 8.747092 10.367287 8.328743 13.190562 10.659016
```

```
set.seed(1)
rnorm(5, mean = 10, sd = 2)
```

```
[1] 8.747092 10.367287 8.328743 13.190562 10.659016
```

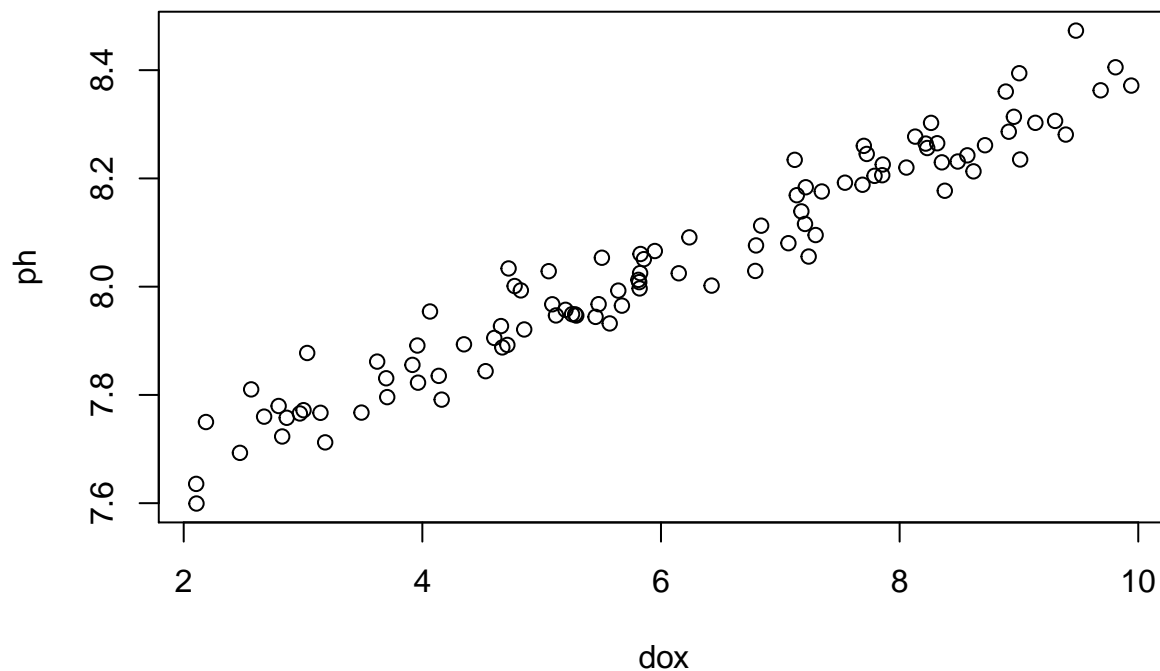
```
# choose a different random 5
rnorm(5, mean = 10, sd = 2)
```

```
[1] 8.359063 10.974858 11.476649 11.151563 9.389223
```

Simulated data

We can use these random number generators to simulate a linear relationship between dissolved oxygen and pH that we can then model:

```
dox <- runif(100, 2, 10)
ph <- 7.5 + (dox * 0.09) + rnorm(100, 0, 0.05)
plot(dox, ph)
```



Linear models

In R, models are usually based on `formula` objects. Formulae are constructed using the tilde (`~`) operator. The syntax is `y ~ x`, which is translated as `y` is a function of `x`. Here's an example for our pH and dissolved oxygen simulated data:

```
pd.form <- ph ~ dox
str(pd.form)
```

```
Class 'formula' language ph ~ dox
..- attr(*, ".Environment")=<environment: R_GlobalEnv>
```

The function `lm` fits a linear model to a formula and returns the intercept and slope estimates as well as diagnostics of the fit:

```
pd.lm <- lm(ph ~ dox)
# Here's a simple summary of the fit
print(pd.lm)
```

Call:

```
lm(formula = ph ~ dox)
```

Coefficients:

```
(Intercept)      dox
  7.49528      0.09031
```

Here are all of the elements in the fitted object:

```
str(pd.lm)
```

List of 12

```
$ coefficients : Named num [1:2] 7.4953 0.0903
..- attr(*, "names")= chr [1:2] "(Intercept)" "dox"
$ residuals    : Named num [1:100] 0.12187 0.00161 0.03698 0.00519 -0.03372 ...
..- attr(*, "names")= chr [1:100] "1" "2" "3" "4" ...
$ effects      : Named num [1:100] -80.40688 1.93363 0.02556 -0.00457 -0.04393 ...
..- attr(*, "names")= chr [1:100] "(Intercept)" "dox" "" "" ...
$ rank         : int 2
$ fitted.values: Named num [1:100] 8.35 7.83 8.15 7.77 7.87 ...
..- attr(*, "names")= chr [1:100] "1" "2" "3" "4" ...
$ assign       : int [1:2] 0 1
$ qr           :List of 5
..$ qr        : num [1:100, 1:2] -10 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 0.1 ...
.. ..- attr(*, "dimnames")=List of 2
.. .. ..$ : chr [1:100] "1" "2" "3" "4" ...
.. .. ..$ : chr [1:2] "(Intercept)" "dox"
.. ..- attr(*, "assign")= int [1:2] 0 1
..$ qraux: num [1:2] 1.1 1.12
..$ pivot: int [1:2] 1 2
..$ tol   : num 1e-07
..$ rank  : int 2
..- attr(*, "class")= chr "qr"
$ df.residual : int 98
$ xlevels     : Named list()
$ call        : language lm(formula = ph ~ dox)
$ terms       :Classes 'terms', 'formula' language ph ~ dox
.. ..- attr(*, "variables")= language list(ph, dox)
.. ..- attr(*, "factors")= int [1:2, 1] 0 1
.. .. ..- attr(*, "dimnames")=List of 2
.. .. .. ..$ : chr [1:2] "ph" "dox"
.. .. .. ..$ : chr "dox"
.. ..- attr(*, "term.labels")= chr "dox"
.. ..- attr(*, "order")= int 1
.. ..- attr(*, "intercept")= int 1
.. ..- attr(*, "response")= int 1
.. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
.. ..- attr(*, "predvars")= language list(ph, dox)
.. ..- attr(*, "dataClasses")= Named chr [1:2] "numeric" "numeric"
.. .. ..- attr(*, "names")= chr [1:2] "ph" "dox"
```

```

$ model      : 'data.frame': 100 obs. of 2 variables:
..$ ph : num [1:100] 8.47 7.83 8.18 7.77 7.84 ...
..$ dox: num [1:100] 9.48 3.7 7.21 3 4.14 ...
..- attr(*, "terms")=Classes 'terms', 'formula' language ph ~ dox
.. . . .- attr(*, "variables")= language list(ph, dox)
.. . . .- attr(*, "factors")= int [1:2, 1] 0 1
.. . . .- attr(*, "dimnames")=List of 2
.. . . . $ : chr [1:2] "ph" "dox"
.. . . . $ : chr "dox"
.. . . .- attr(*, "term.labels")= chr "dox"
.. . . .- attr(*, "order")= int 1
.. . . .- attr(*, "intercept")= int 1
.. . . .- attr(*, "response")= int 1
.. . . .- attr(*, ".Environment")=<environment: R_GlobalEnv>
.. . . .- attr(*, "predvars")= language list(ph, dox)
.. . . .- attr(*, "dataClasses")= Named chr [1:2] "numeric" "numeric"
.. . . .- attr(*, "names")= chr [1:2] "ph" "dox"
- attr(*, "class")= chr "lm"

```

The estimated coefficients are stored in the `$coefficients` element:

```
pd.lm$coefficients
```

```

(Intercept)      dox
  7.4952814    0.0903088

```

They can also be obtained with a call to the `coef()` function:

```
coef(pd.lm)
```

```

(Intercept)      dox
  7.4952814    0.0903088


```

We can use the `abline` function to plot the estimated fit over the data:

```

plot(dox, ph)
abline(pd.lm, col = "red", lwd = 2, lty = "dashed")

```



More detail about the fit can be extracted with the `summary` function. In particular, we can see a summary of the residuals to inspect normality of the errors, as well as the standard errors and p-values for tests of significant deviation of the estimated parameters from zero:

```
x <- summary(pd.lm)
print(x)
```

Call:

```
lm(formula = ph ~ dox)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.093234	-0.029194	-0.004968	0.027440	0.121873

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	7.495281	0.014262	525.53	<2e-16 ***
dox	0.090309	0.002226	40.57	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.04766 on 98 degrees of freedom

Multiple R-squared: 0.9438, Adjusted R-squared: 0.9432

F-statistic: 1646 on 1 and 98 DF, p-value: < 2.2e-16

```
str(x)
```

List of 11

```
$ call      : language lm(formula = ph ~ dox)
$ terms     :Classes 'terms', 'formula' language ph ~ dox
.. ..- attr(*, "variables")= language list(ph, dox)
.. ..- attr(*, "factors")= int [1:2, 1] 0 1
.. ..- attr(*, "dimnames")=List of 2
```

```

.. ..$ : chr [1:2] "ph" "dox"
.. ..$ : chr "dox"
.. ..- attr(*, "term.labels")= chr "dox"
.. ..- attr(*, "order")= int 1
.. ..- attr(*, "intercept")= int 1
.. ..- attr(*, "response")= int 1
.. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
.. ..- attr(*, "predvars")= language list(ph, dox)
.. ..- attr(*, "dataClasses")= Named chr [1:2] "numeric" "numeric"
.. ..- attr(*, "names")= chr [1:2] "ph" "dox"
$ residuals : Named num [1:100] 0.12187 0.00161 0.03698 0.00519 -0.03372 ...
..- attr(*, "names")= chr [1:100] "1" "2" "3" "4" ...
$ coefficients : num [1:2, 1:4] 7.50 9.03e-02 1.43e-02 2.23e-03 5.26e+02 ...
..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:2] "(Intercept)" "dox"
.. ..$ : chr [1:4] "Estimate" "Std. Error" "t value" "Pr(>|t|)"
$ aliased : Named logi [1:2] FALSE FALSE
..- attr(*, "names")= chr [1:2] "(Intercept)" "dox"
$ sigma : num 0.0477
$ df : int [1:3] 2 98 2
$ r.squared : num 0.944
$ adj.r.squared: num 0.943
$ fstatistic : Named num [1:3] 1646 1 98
..- attr(*, "names")= chr [1:3] "value" "numdf" "dendf"
$ cov.unscaled : num [1:2, 1:2] 0.08956 -0.01317 -0.01317 0.00218
..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:2] "(Intercept)" "dox"
.. ..$ : chr [1:2] "(Intercept)" "dox"
- attr(*, "class")= chr "summary.lm"

```

We can use the model fit object to predict new data too. We just need a data frame of the new values with column names of the independent values the same as those in the original model:

```

new.dox <- data.frame(dox = runif(5, 2, 10))
predict(pd.lm, new.dox)

      1      2      3      4      5
8.156066 7.901664 7.968998 8.395535 8.293671

# include confidence intervals
predict(pd.lm, new.dox, interval = "confidence")

```

```

      fit      lwr      upr
1 8.156066 8.145053 8.167079
2 7.901664 7.890015 7.913312
3 7.968998 7.958911 7.979084
4 8.395535 8.375770 8.415301
5 8.293671 8.278097 8.309245

```

Models can be built on categorical predictors as well. In this example we test whether or not surface temperature differs among the first five stations:

```

ctd <- read.csv("ctd.csv", stringsAsFactors = FALSE)
surf <- subset(ctd, depth == 1 & station %in% paste0("Station.", 1:5))

temp.lm <- lm(temp ~ station, surf)
summary(temp.lm)

```

```

Call:
lm(formula = temp ~ station, data = surf)

Residuals:
    Min       1Q   Median       3Q      Max
-4.583 -1.838 -0.160  1.367  6.331

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    17.2263     0.3090  55.749 < 2e-16 ***
stationStation.2 -0.3656     0.4370  -0.837  0.40349
stationStation.3 -0.4963     0.4370  -1.136  0.25703
stationStation.4 -0.8671     0.4370  -1.984  0.04816 *
stationStation.5 -1.2236     0.4370  -2.800  0.00545 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.373 on 290 degrees of freedom
Multiple R-squared:  0.03099,    Adjusted R-squared:  0.01762
F-statistic: 2.318 on 4 and 290 DF,  p-value: 0.05723

```

Note that the results list the dummy variables representing the levels of the categorical station predictor. Their estimated effects are expressed as being relative to the first level.

ANOVA

An analysis of variance (ANOVA) is similar to the multi-category linear model and gets specified with the same formula using the `aov` function:

```

temp.aov <- aov(temp ~ station, surf)
summary(temp.aov)

```

```

              Df Sum Sq Mean Sq F value Pr(>F)
station         4   52.2   13.060    2.318 0.0572 .
Residuals      290 1633.6    5.633
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
str(temp.aov)
```

List of 13

```

$ coefficients : Named num [1:5] 17.226 -0.366 -0.496 -0.867 -1.224
..- attr(*, "names")= chr [1:5] "(Intercept)" "stationStation.2" "stationStation.3" "stationStation.4"
$ residuals    : Named num [1:295] 3.09 -2.99 -1.52 -3.03 -2.74 ...
..- attr(*, "names")= chr [1:295] "18" "40" "73" "103" ...
$ effects      : Named num [1:295] -285.73 1.93 1.33 -1.6 -6.65 ...
..- attr(*, "names")= chr [1:295] "(Intercept)" "stationStation.2" "stationStation.3" "stationStation.4"
$ rank         : int 5
$ fitted.values: Named num [1:295] 17.2 17.2 17.2 17.2 17.2 ...
..- attr(*, "names")= chr [1:295] "18" "40" "73" "103" ...
$ assign       : int [1:5] 0 1 1 1 1
$ qr           :List of 5
..$ qr        : num [1:295, 1:5] -17.1756 0.0582 0.0582 0.0582 0.0582 ...
.. ..- attr(*, "dimnames")=List of 2

```

```

.. .. ..$ : chr [1:295] "18" "40" "73" "103" ...
.. .. ..$ : chr [1:5] "(Intercept)" "stationStation.2" "stationStation.3" "stationStation.4" ...
.. ..- attr(*, "assign")= int [1:5] 0 1 1 1 1
.. ..- attr(*, "contrasts")=List of 1
.. .. ..$ station: chr "contr.treatment"
..$ graux: num [1:5] 1.06 1.03 1.03 1.05 1.08
..$ pivot: int [1:5] 1 2 3 4 5
..$ tol : num 1e-07
..$ rank : int 5
..- attr(*, "class")= chr "qr"
$ df.residual : int 290
$ contrasts :List of 1
..$ station: chr "contr.treatment"
$ xlevels :List of 1
..$ station: chr [1:5] "Station.1" "Station.2" "Station.3" "Station.4" ...
$ call : language aov(formula = temp ~ station, data = surf)
$ terms :Classes 'terms', 'formula' language temp ~ station
.. ..- attr(*, "variables")= language list(temp, station)
.. ..- attr(*, "factors")= int [1:2, 1] 0 1
.. .. ..- attr(*, "dimnames")=List of 2
.. .. .. ..$ : chr [1:2] "temp" "station"
.. .. .. ..$ : chr "station"
.. ..- attr(*, "term.labels")= chr "station"
.. ..- attr(*, "order")= int 1
.. ..- attr(*, "intercept")= int 1
.. ..- attr(*, "response")= int 1
.. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
.. ..- attr(*, "predvars")= language list(temp, station)
.. ..- attr(*, "dataClasses")= Named chr [1:2] "numeric" "character"
.. .. ..- attr(*, "names")= chr [1:2] "temp" "station"
$ model :'data.frame': 295 obs. of 2 variables:
..$ temp : num [1:295] 20.3 14.2 15.7 14.2 14.5 ...
..$ station: chr [1:295] "Station.1" "Station.1" "Station.1" "Station.1" ...
..- attr(*, "terms")=Classes 'terms', 'formula' language temp ~ station
.. .. ..- attr(*, "variables")= language list(temp, station)
.. .. ..- attr(*, "factors")= int [1:2, 1] 0 1
.. .. .. ..- attr(*, "dimnames")=List of 2
.. .. .. .. ..$ : chr [1:2] "temp" "station"
.. .. .. .. ..$ : chr "station"
.. .. ..- attr(*, "term.labels")= chr "station"
.. .. ..- attr(*, "order")= int 1
.. .. ..- attr(*, "intercept")= int 1
.. .. ..- attr(*, "response")= int 1
.. .. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
.. .. ..- attr(*, "predvars")= language list(temp, station)
.. .. ..- attr(*, "dataClasses")= Named chr [1:2] "numeric" "character"
.. .. .. ..- attr(*, "names")= chr [1:2] "temp" "station"
- attr(*, "class")= chr [1:2] "aov" "lm"

```

The analysis of variance table can also be computed from an `lm` object using `anova`:

```
anova(temp.lm)
```

Analysis of Variance Table


```

Response: temp
      Df Sum Sq Mean Sq F value Pr(>F)
station  4  52.24  13.0602   2.3184 0.05723 .
Residuals 290 1633.62   5.6332
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Differences between levels of the predictor can be tested with the TukeyHSD() function:
TukeyHSD(temp.aov)

```

```

      Tukey multiple comparisons of means
      95% family-wise confidence level

Fit: aov(formula = temp ~ station, data = surf)

$station
      diff      lwr      upr      p adj
Station.2-Station.1 -0.3655932 -1.565124  0.83393726 0.9190166
Station.3-Station.1 -0.4962712 -1.695802  0.70325929 0.7874740
Station.4-Station.1 -0.8671186 -2.066649  0.33241184 0.2762382
Station.5-Station.1 -1.2235593 -2.423090 -0.02402884 0.0430545
Station.3-Station.2 -0.1306780 -1.330208  1.06885251 0.9982509
Station.4-Station.2 -0.5015254 -1.701056  0.69800506 0.7809047
Station.5-Station.2 -0.8579661 -2.057497  0.34156438 0.2866747
Station.4-Station.3 -0.3708475 -1.570378  0.82868302 0.9150329
Station.5-Station.3 -0.7272881 -1.926819  0.47224234 0.4577427
Station.5-Station.4 -0.3564407 -1.555971  0.84308980 0.9256832

```

Non-linear models

To illustrate non-linear model fitting and statistical tests, we'll first create a function to simulate growth data (length ~ age) based on a Gompertz curve. The Gompertz function is $length = L_0 \cdot e^{k(1 - e^{-g \cdot age})}$, where L_0 is the length at birth (LAB). Here's the function to create simulated growth data:

```

# age.range - a two element vector giving the minimum and maximum ages
# lab - the length at birth
# k, g - displacement and rate parameters
# std.dev - standard deviation for the error term
# sample.size - number of points to simulate


sim.growth.func <- function(age.range, lab, k, g,
  std.dev, sample.size) {
  # Check to make sure age.range is a reasonable vector
  if (!is.numeric(age.range) || !is.vector(age.range))
    stop("'age.range' is not a numeric vector")
  if (any(age.range < 0)) stop("'age.range' < 0")
  if (age.range[1] >= age.range[2])
    stop("'age.range[1]' >= 'age.range[2]'"")
  # Generate some random ages between min and max of age.range
  ages <- runif(sample.size, age.range[1], age.range[2])
  # Calculate the expected length for those ages from the Gompertz equation
  expected.length <- lab * exp(k * (1 - exp(-g * ages)))
  # Add some error to the lengths and return the named array
  length.err <- rnorm(sample.size, 0, std.dev)

```

```
as.data.frame(cbind(age = ages, length = expected.length + length.err))
}
```

With this function, we can now simulate some growth data:

```
growth.df <- sim.growth.func(
  age.range = c(0, 65),
  lab = 10,
  k = 2,
  g = 0.25,
  std.dev = 5,
  sample.size = 200
)
plot(length ~ age, growth.df)
```



Now let's use nonlinear least squares to estimate the parameters from this simulated data. We can do that with the `nls` function, which behaves very similarly to `lm`. The main difference is that we need to supply initial values, which are specified in the third argument, `start`, which should be chosen carefully so as to ensure convergence.

```
gr.form <- length ~ lab * exp(k * (1 - exp(-g * age)))
# starting values for k and g are too far off for default number of iterations
gr.nls <- nls(gr.form, growth.df, start = c(lab = 15, k = 10, g = 10))
```

Error in `numericDeriv(form[[3L]], names(ind), env)`: Missing value or an infinity produced when evaluating

```
# this should work
gr.nls <- nls(gr.form, growth.df, start = c(lab = 15, k = 1, g = 0.5))
print(gr.nls)
```

Nonlinear regression model

```
model: length ~ lab * exp(k * (1 - exp(-g * age)))
data: growth.df
lab      k      g
8.309 2.182 0.272
```

residual sum-of-squares: 5498

Number of iterations to convergence: 6
Achieved convergence tolerance: 2.8e-06

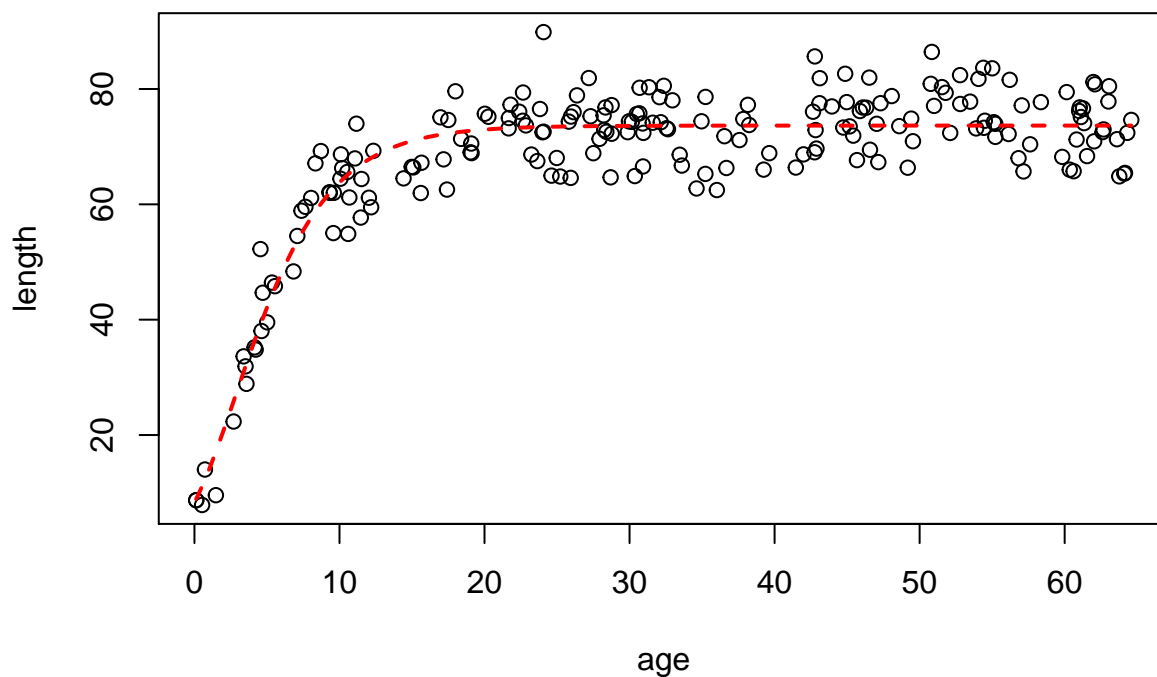
...and here are the estimated coefficients:

```
gr.coef <- coef(gr.nls)
gr.coef
```

```
      lab      k      g
8.3093886 2.1819649 0.2720253
```

We'll plot the fitted curve:

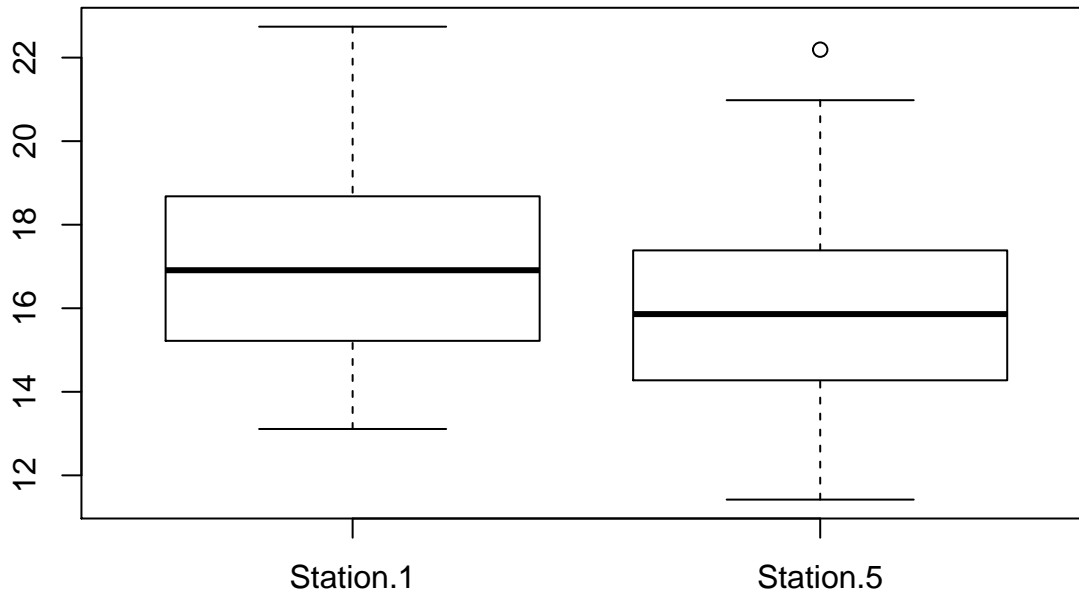
```
grow.fit <- data.frame(
  age = seq(
    min(growth.df$age),
    max(growth.df$age),
    length.out = 1000
  )
)
grow.fit$length <- predict(gr.nls, grow.fit)
plot(length ~ age, growth.df)
lines(grow.fit$age, grow.fit$length, col = "red", lwd = 2, lty = "dashed")
```



Statistical tests

There are several functions for standard statistical tests that all have similar outputs. The most common ones are `binom.test`, `chisq.test`, `kruskal.test`, `ks.test`, and `t.test`. As an example, we'll conduct a t-test of surface temperatures between stations 1 and 5.

```
surf15 <- subset(surf, station %in% c("Station.1", "Station.5"))
boxplot(temp ~ station, surf15)
```



```
surf15.t <- t.test(temp ~ station, surf15)
surf15.t
```

Welch Two Sample t-test

```
data: temp by station
t = 2.7601, df = 115.31, p-value = 0.006725
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 0.3454737 2.1016449
sample estimates:
mean in group Station.1 mean in group Station.5
      17.22627           16.00271
```

```
str(surf15.t)
```

```
List of 9
 $ statistic : Named num 2.76
  ..- attr(*, "names")= chr "t"
 $ parameter : Named num 115
  ..- attr(*, "names")= chr "df"
 $ p.value    : num 0.00672
 $ conf.int   : atomic [1:2] 0.345 2.102
  ..- attr(*, "conf.level")= num 0.95
 $ estimate   : Named num [1:2] 17.2 16
  ..- attr(*, "names")= chr [1:2] "mean in group Station.1" "mean in group Station.5"
 $ null.value : Named num 0
  ..- attr(*, "names")= chr "difference in means"
 $ alternative: chr "two.sided"
 $ method     : chr "Welch Two Sample t-test"
 $ data.name  : chr "temp by station"
 - attr(*, "class")= chr "htest"
```

For a chi-squared test, there's the function `chisq.test`. As an example, we'll see if ctd sampling is random across years and months. Let's first examine years, using the output from `table`:

```
# Extract year, month, and day as columns:
ctd$year <- as.numeric(substr(ctd$sample_date, 1, 4))
ctd$month <- as.numeric(substr(ctd$sample_date, 6, 7))
ctd$day <- as.numeric(substr(ctd$sample_date, 9, 10))

yr.freq <- table(ctd$year)
yr.chisq <- chisq.test(yr.freq)
str(yr.chisq)
```

```
List of 9
 $ statistic: Named num 5658
  ..- attr(*, "names")= chr "X-squared"
 $ parameter: Named num 6
  ..- attr(*, "names")= chr "df"
 $ p.value   : num 0
 $ method    : chr "Chi-squared test for given probabilities"
 $ data.name : chr "yr.freq"
 $ observed  : 'table' int [1:7(1d)] 12886 13908 13865 13907 6760 8195 8120
  ..- attr(*, "dimnames")=List of 1
  .. ..$ : chr [1:7] "2010" "2011" "2012" "2013" ...
 $ expected  : Named num [1:7] 11092 11092 11092 11092 11092 ...
  ..- attr(*, "names")= chr [1:7] "2010" "2011" "2012" "2013" ...
 $ residuals: table [1:7(1d)] 17 26.7 26.3 26.7 -41.1 ...
  ..- attr(*, "dimnames")=List of 1
  .. ..$ : chr [1:7] "2010" "2011" "2012" "2013" ...
 $ stdres    : table [1:7(1d)] 18.4 28.9 28.4 28.9 -44.4 ...
  ..- attr(*, "dimnames")=List of 1
  .. ..$ : chr [1:7] "2010" "2011" "2012" "2013" ...
 - attr(*, "class")= chr "htest"
```

Here's a test of sampling frequency by year and month:

```
yr.mo.chisq <- chisq.test(table(ctd$month, ctd$year))
yr.mo.chisq
```

Pearson's Chi-squared test

```
data: table(ctd$month, ctd$year)
X-squared = 7509.3, df = 66, p-value < 2.2e-16
```

```
str(yr.mo.chisq)
```

```
List of 9
 $ statistic: Named num 7509
  ..- attr(*, "names")= chr "X-squared"
 $ parameter: Named int 66
  ..- attr(*, "names")= chr "df"
 $ p.value   : num 0
 $ method    : chr "Pearson's Chi-squared test"
 $ data.name : chr "table(ctd$month, ctd$year)"
 $ observed  : 'table' int [1:12, 1:7] 1154 1136 1157 179 1153 1147 1157 1158 1164 1163 ...
  ..- attr(*, "dimnames")=List of 2
  .. ..$ : chr [1:12] "1" "2" "3" "4" ...
  .. ..$ : chr [1:7] "2010" "2011" "2012" "2013" ...
 $ expected  : num [1:12, 1:7] 919 1390 926 775 1393 ...
```

```

..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:12] "1" "2" "3" "4" ...
.. ..$ : chr [1:7] "2010" "2011" "2012" "2013" ...
$ residuals: table [1:12, 1:7] 7.77 -6.81 7.59 -21.42 -6.42 ...
..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:12] "1" "2" "3" "4" ...
.. ..$ : chr [1:7] "2010" "2011" "2012" "2013" ...
$ stdres : table [1:12, 1:7] 8.82 -7.89 8.62 -24.19 -7.45 ...
..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:12] "1" "2" "3" "4" ...
.. ..$ : chr [1:7] "2010" "2011" "2012" "2013" ...
- attr(*, "class")= chr "htest"

```

...we can get also the same result by setting the x and y arguments of `chisq.test` to the original columns:

```
chisq.test(ctd$month, ctd$year)
```

Pearson's Chi-squared test

```
data: ctd$month and ctd$year
X-squared = 7509.3, df = 66, p-value < 2.2e-16
```

The `chisq.test` function also has the ability to estimate significance via a bootstrap, which is selected by setting `simulate.p.value = TRUE`:

```
chisq.test(ctd$month, ctd$year, simulate.p.value = TRUE)
```

Pearson's Chi-squared test with simulated p-value (based on 2000 replicates)

```
data: ctd$month and ctd$year
X-squared = 7509.3, df = NA, p-value = 0.0004998
```

Bootstrapping

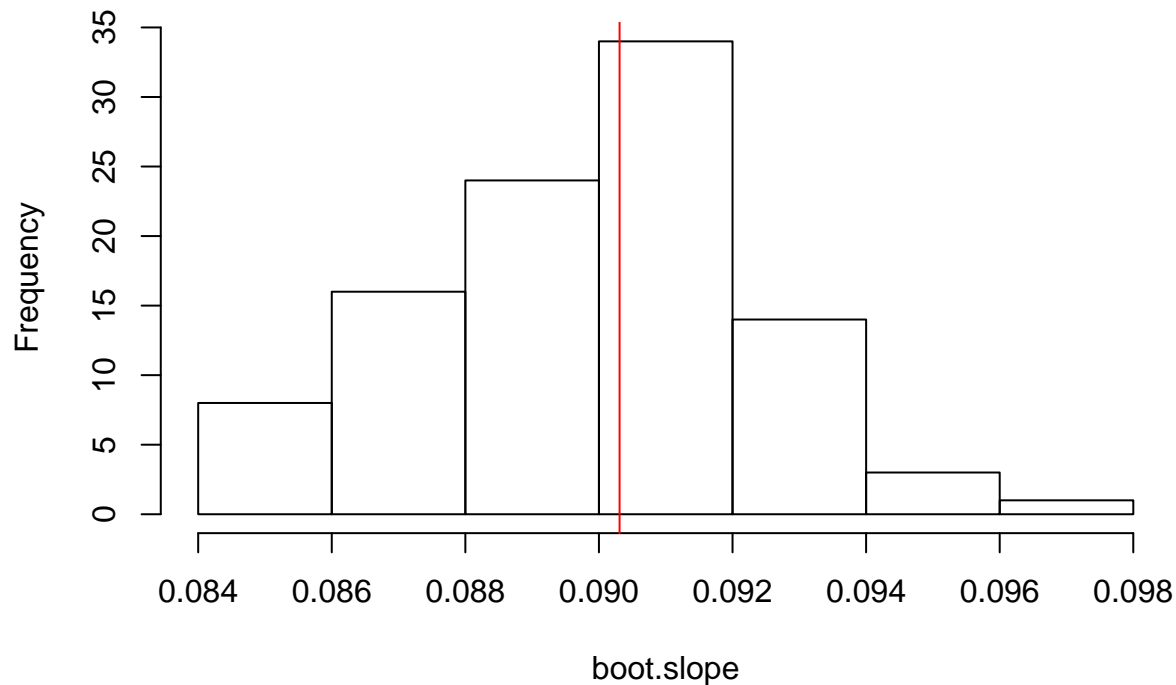
We can get bootstrap estimates of the variance of a parameter by randomly selecting rows of our original data with replacement. As an example, here's some code for getting the variance of the slope parameter from our simulated ph/dox data from 100 bootstrap samples:

```

# calculate the observed value
obs.slope <- coef(lm(ph ~ dox))[2]
# conduct 100 bootstrap replicates
boot.slope <- sapply(1:100, function(i) {
  boot.i <- sample(1:length(ph), length(ph), replace = TRUE)
  boot.ph <- ph[boot.i]
  boot.dox <- dox[boot.i]
  boot.lm <- lm(boot.ph ~ boot.dox)
  coef(boot.lm)[2]
})
# a histogram of the bootstrap slopes with the observed value in red
hist(boot.slope)
abline(v = obs.slope, col = "red")

```

Histogram of boot.slope



```
# the observed value and a summary of the bootstrap replicates
obs.slope
```

```
dox
0.0903088
```

```
summary(boot.slope)
```

```
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.08427 0.08836 0.09017 0.08992 0.09170 0.09655
```

Permutation tests

We can also use resampling to conduct permutation tests where we create a null distribution by randomly permuting the predictor variable. Here we simulate separate male and female growth and test if adults (> 18 y/o) have different lengths.

```
# Simulate data for the two sexes
m.growth <- sim.growth.func(c(0, 65), 10, 2.05, 0.27, 5, 100)
f.growth <- sim.growth.func(c(0, 65), 10, 2.02, 0.25, 4, 100)
# Extract adults
adult.m <- m.growth[m.growth[, "age"] > 18, ]
adult.f <- f.growth[f.growth[, "age"] > 18, ]

# Create combined sex data.frame
adult.df <- rbind(adult.m, adult.f)
adult.df$sex <- rep(c("m", "f"), c(nrow(adult.m), nrow(adult.f)))

# the t-test result
t.test(length ~ sex, adult.df)
```

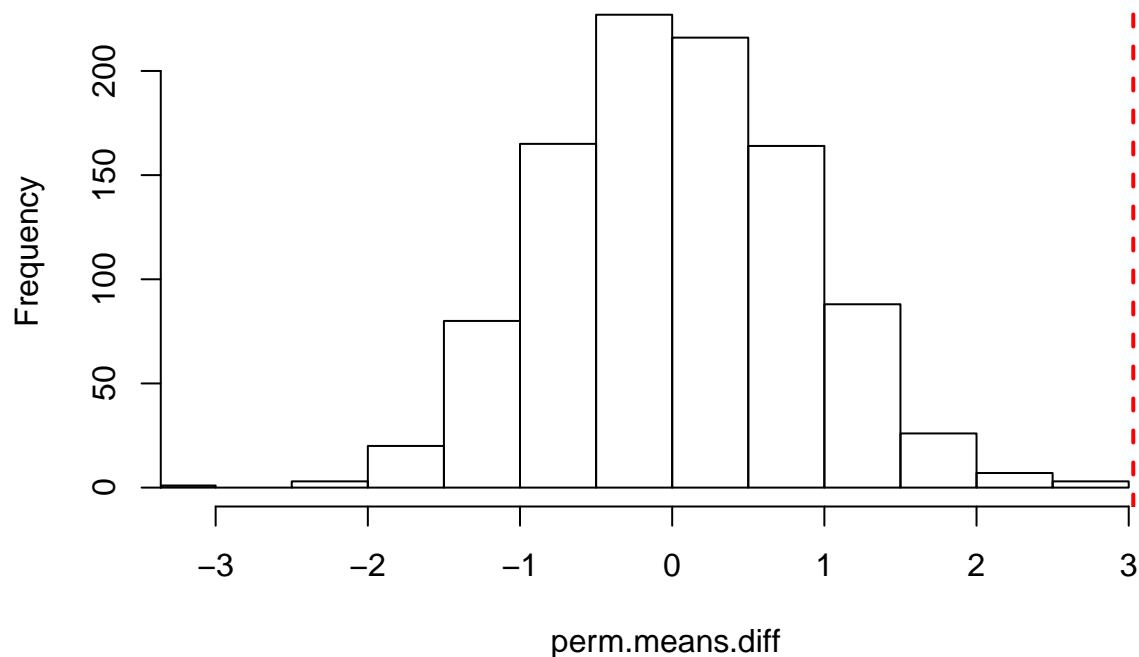
Welch Two Sample t-test

```
data: length by sex
t = -3.9812, df = 137.6, p-value = 0.0001105
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -4.534991 -1.525090
sample estimates:
mean in group f mean in group m
 74.62510      77.65514
```

```
# calculate observed difference of means
obs.means <- tapply(adult.df$length, adult.df$sex, mean)
obs.means.diff <- diff(obs.means)
# permute sexes and calculate difference of means for null distribution
perm.means <- sapply(1:1000, function(i) {
  tapply(adult.df$length, sample(adult.df$sex), mean)
})
perm.means.diff <- apply(perm.means, 2, diff)

# plot histogram of permutation differences with observed difference
hist(perm.means.diff, xlim = range(c(obs.means.diff, perm.means.diff)))
abline(v = obs.means.diff, col = "red", lty = "dashed", lwd = 2)
```

Histogram of perm.means.diff



```
# observed difference
obs.means.diff
```

```
m
3.030041
```



```
# summary of permutation differences
summary(perm.means.diff)

      Min.    1st Qu.     Median       Mean    3rd Qu.      Max.
-3.115000 -0.535300  0.008564  0.021450  0.595000  2.574000

# fraction of permutation >= observed
mean(perm.means.diff >= obs.means.diff)

[1] 0
```